

US market analysis — AI-spend attribution & governance (Outlay)

Question: Who buys AI in the US, which variables predict they get real value from Outlay (metadata-only, BYOK, read-only attribution + forecasting + budget governance of AI spend), how do they buy AI compute today, and what's the defensible TAM — core vs incremental?

Provenance. Built from a deep-research run (fan-out web search → fetch → **3-vote adversarial verification**). 19 claims survived verification; 5 were refuted and excluded (listed at the end so we don't repeat them). Macro facts below are **cited**; the taxonomy, fit-scoring, purchasing breakdown, and TAM are **Outlay's analysis on top of those facts**, with assumptions labeled. Scope: **US**, demand-side AI spend (what companies pay to use AI — SaaS + API + their own cloud AI), which is the pool Outlay can govern. (Supply-side data-center capex is sized separately and is **not** Outlay's addressable pool — see TAM caveat.)

0. Executive summary

1. **The category is real and inflecting now.** AI-spend management went from a niche to near- universal in 24 months — **98% of organizations now manage AI spend as part of FinOps, up from 31% in 2024**, the fastest cost-discipline adoption the FinOps survey has ever recorded, and it's the **#1 FinOps priority** for the next 12 months. [FinOps Foundation 2026]
 2. **The pain is structural, not cyclical.** AI bills scale by **tokens/compute, not seats**, so per-seat pricing can't price them; spend is **consumption-based, opaque, account-level, and spread across many providers with no unified view** — the exact attribution gap Outlay closes. [Flexera; Datadog; Vantage]
 3. **The bill is large, growing, and finance-visible.** **88% of large US enterprises plan to increase GenAI budgets in the next 12 months; 62% by ≥10%.** Enterprise LLM **API** spend alone was **~\$8.4B** by mid-2025 (≈2× in ~6 months), and the field is **multi-vendor** (Anthropic ~32–40%, OpenAI ~25–27%, Google ~20%, Llama ~9%). [Wharton/GBK; Menlo Ventures]
 4. **The wedge is the engineering/coding-agent segment.** Code generation is AI's first breakout use case (**~\$4B, ~28% of enterprise LLM deployments**, Claude ~42% developer share), it runs on **raw API tokens + coding agents** (the worst attribution gap), and **inference/production has overtaken training** — i.e., recurring usage cost, not one-time. [Menlo Ventures]
 5. **TAM (Outlay model).** **US governable AI-spend pool ≈ \$35–70B in 2026**, growing ~50–100%/yr. At a FinOps-style **2–5% software take-rate**, that's a **~\$1–3B US revenue TAM**, of which the **core product** (attribution + forecast + budget governance) is the entry wedge and **incremental features** (enforcement, chargeback, procurement optimization) roughly double the wallet over time.
-

1. Market context (the facts the strategy rests on)

AI-cost governance has become a board-level FinOps discipline — fast. - **98% of organizations now manage AI spend** as part of FinOps, **up from 31% two years ago** (2024); the survey calls it the fastest technology-cost adoption in its history. [FinOps Foundation, *State of FinOps 2026*, n=1,192 orgs representing \$83B+ tech spend — data.finops.org] (*high confidence*) - **Managing AI spend is the #1 FinOps priority** for the next 12 months (cited by 32.7% of teams — a plurality), and **AI cost management is the #1 skillset** teams want to build (58%). [FinOps

2026] (high) - FinOps scope is expanding beyond cloud: **90% manage (or plan to manage) SaaS, 64% licensing, 57% private cloud, 48% data center.** [FinOps 2026] (high) — i.e., the buyer Outlay sells to already exists and is actively widening their remit to exactly this.

The bill is big, growing, multi-vendor, and consumption-shaped. - 88% of US enterprises (1,000+ employees, >\$50M revenue) plan to increase GenAI budgets in the next 12 months; 62% expect ≥10% increases.

[Wharton/GBK Collective, 2025 AI Adoption Report, ~800 US enterprise decision-makers — ai.wharton.upenn.edu] (high) - Enterprise **LLM API spend ≈ \$8.4B by mid-2025**, roughly double ~6 months prior; basis is a Menlo survey of 150 technical leaders. [Menlo Ventures, 2025 Mid-Year LLM Market Update] (medium — the "\$8.4B" figure is corroborated; the specific "doubled in exactly 6 months" framing failed verification) - **Multi-vendor field:** Anthropic ~32% (year-end ~40%), OpenAI ~25% (down from ~50% in 2023), Google ~20%, Meta Llama ~9%. [Menlo] (high on the numbers; nuance: top-3 ≈ 88% concentration and only ~11% of teams switched providers in the past year — vendor diversity is real at the market level and growing within large enterprises, but isn't yet universal per-company multi-homing) - **Inference (production) has overtaken training** — 74% of startups and 49% of large enterprises say most compute is now inference — i.e., **recurring usage cost**, not one-time training. [Menlo] (high) - **Code generation is the breakout use case** (~\$4B, ~28% of enterprise LLM deployments; Claude ~42% developer share vs OpenAI ~21%). [Menlo] (high)

Why per-seat pricing broke (the structural opening). - AI features scale by tokens, compute-minutes, and model complexity — not by user — so per-seat pricing structurally can't price them; this is the core driver of the **seat → consumption** shift (~65% of vendors now use hybrid seat+usage pricing). [Flexera; corroborated by Bain] (high) - **Consumption data is opaque and aggregated at the account level**, and AI usage is **spread across internal platforms, SaaS tools, and vendor APIs with no out-of-the-box unified view** — making attribution to workloads/teams hard. Provider billing "typically breaks down costs only to the model, project, or workspace level," and "team/cost-center metadata doesn't exist natively in AI billing data." [Flexera; Datadog; Vantage] (high) — **this is Outlay's problem statement, verbatim. - Finance often can't map a credit back to a workload**; credit expirations/resets/rollover "silently create overages or strand value," producing budget surprises and uncontrolled use. A reported **78% of IT leaders saw surprise AI charges** in the past year. [Flexera] (high) - **Seat overbuying already wastes money** (~51–53% of SaaS licenses go unused; large enterprises waste ~\$21M/yr), and AI makes the meter worse. [Flexera; CFO Dive] (high)

Supply-side context (NOT the addressable pool). Global **AI data-center capex is forecast at \$400–450B in 2026** (\$250–300B of it chips), rising to **~\$1T by 2028**; inference ≈ two-thirds of all compute in 2026. [Deloitte, TMT Predictions 2026] (high). This is the hyperscaler/AI-lab **build**; Outlay governs the **consuming company's** AI bill, so capex is the backdrop, not the TAM (see §4 caveat).

2. Customer taxonomy — the variables, and what each implies for fit

Categorize US AI buyers on these axes. For each, the **"hot" end** (higher Outlay value) is bolded.

Variable	Buckets (low-fit → high-fit)	Why it predicts value
AI spend model	seat-SaaS only → mixed → heavy API/ consumption + credits	Attribution gap lives in token/usage spend, not flat seats [Flexera]
Spend magnitude & growth	flat/small → large & growing ≥10%/yr	88% increasing budgets; bigger bill = bigger problem [Wharton]
# providers / tools		Fragmentation = no native rollup [Menlo; Datadog]

Variable	Buckets (low-fit → high-fit)	Why it predicts value
	single tool → many tools, multi-provider, no unified view	
Cost-owner gap	one owner, in control → no owner / finance blind	"can't map credit → workload" [Flexera]
AI maturity / current state	laggard → experimenter → scaled adopter (production inference) → AI-native	Inference = recurring cost; experiments are too small to govern [Menlo]
Employee mix	non-technical, light AI → engineering/data-heavy + AI-native knowledge work	Eng on coding agents = highest token spend [Menlo]
Size (employees)	<50 → 50–250 → 250–5k → 5k+	Big enough to have a real bill + departments to attribute; not so big they've already built it
Stage / type	pre-seed → growth/late-stage startup, PE-backed, established corp digitizing → public mega-cap	Budget pressure + finance scrutiny + can't-yet-build-it-themselves
Vertical	low-AI (e.g. traditional mfg) → software/tech, financial services, digital-native, prof. services, gov/regulated	Where AI spend concentrates and where governance is mandated
Key tools	one chat seat → Cursor/Copilot + Claude/OpenAI/Bedrock APIs + Copilot/ChatGPT-Ent seats	Coding agents + provider APIs are the consumption core

AI usage by employee function (where the spend actually is): - Engineering / data-ML — highest \$/head. Coding agents (Cursor, Copilot, Claude Code) + raw provider API tokens (OpenAI/Anthropic/Bedrock/Azure). Consumption-based, volatile, hardest to attribute. *This is the core wedge.* [Menlo: code gen \$4B/28%] - **Customer support / service** — high and rising (deflection bots, agent assist); +16pp perceived impact vs 2024. Often usage-metered. [Wharton] - **Sales & marketing** — broad seat adoption (Copilot, ChatGPT, content tools); mostly seats today. - **Legal, Procurement, Finance, HR, Ops** — fastest-rising *perceived* impact (Legal **+24pp**, Procurement **+15pp**) — AI spend is **spreading across departments**, multiplying who needs attribution. [Wharton] Mostly seats now, drifting to usage. - Cross-cutting: **82% of enterprise leaders use GenAI weekly, 46% daily** — adoption is broad-based and recurring (though that's leaders, not yet the whole workforce). [Wharton]

3. The critical value-prop-fit variables (ranked by predictive power)

If you could ask only a few questions to qualify a prospect, ask these — in order:

1. **"How much of your AI spend is consumption/usage-based (API tokens, credits, metered) vs flat seats — and is it growing?"** (*highest signal.*) Heavy + growing consumption spend = the attribution gap is real and worsening. [Flexera; Menlo]
2. **"How many AI tools/providers are in use, and can anyone produce one unified view of the spend and who/what drove it?"** Many tools + "no, nobody can" = textbook Outlay. [Datadog; Vantage]
3. **"Who owns the AI budget, and have you been surprised by an AI bill?"** No clear owner + surprise charges (78% have) = active pain, budget authority to buy. [Flexera]
4. **"What's your AI maturity — experiments, or production workloads running every day?"** Production inference = recurring, governable spend; experiments are too small. [Menlo]
5. **"What share of your headcount is engineering/data, and do they use coding agents / call model APIs directly?"** Eng-heavy = the highest-spend, lowest-visibility function. [Menlo]

6. **Spend magnitude / company size** — gates whether the ROI clears your price (below a threshold the bill is too small to bother governing).

The threshold where the value prop becomes real (Outlay's qualifying line):

A company with **>\$250–500k/yr of AI spend, at least a meaningful share consumption-based, 2+ AI tools/providers, no single owner producing an attributed view, and an engineering/AI-native function** driving usage. Below that, the bill is too small or too simple to need a governance layer (seats with one owner ≈ low fit).

High-fit ICP: growth/late-stage software & digital-native companies (250–5k employees), AI-native or "scaled adopter," engineering-heavy, multi-provider API + coding-agent spend, finance starting to ask "what is this costing per team/project?" Plus **AI-forward financial services and regulated/gov** where governance is *mandated* (your existing Maryland/gov motion is the regulated flavor of this). **Low-fit:** pre-seed/tiny (no bill), AI-laggards (no spend), single-seat-SaaS shops with one owner (little attribution pain — though budget/waste value still exists).

4. How companies buy AI compute today (the central question)

Four models coexist; most companies run **several at once**, which is itself the problem.

Model	Who owns budget	Pricing	Who/what it covers	Why they choose it	Outlay value
(a) Seat / enterprise subscriptions (ChatGPT Enterprise, Copilot, Claude for Work, Glean)	IT / Procurement	per-seat, predictable	knowledge workers (sales, mktg, ops, legal, support)	control, security/SSO, simple budgeting, compliance	budget/showback + waste (overbought seats, ~50% unused)
(b) Central API/ compute + credits/ budgets (OpenAI/ Anthropic/Bedrock/ Azure orgs; cloud GPU)	Engineering / Platform / sometimes Finance	usage/ consumption , volatile	engineers, data/ML, production features, coding agents	speed + developer experience, can't seat-price tokens, scale	core wedge — attribute tokens → team/ticket/ workload, forecast, cap
(c) Reimbursed individual subscriptions	Finance (expensed)	per-seat, scattered	early/SMB, individual power users	speed, low friction, pre-policy	near-zero visibility; consolidation/ showback opportunity
(d) Shadow AI on personal accounts/ cards	nobody (ungoverned)	mixed	anyone, pre-policy	velocity, avoiding procurement	security/compliance + surprise-charge risk; the thing governance teams are chartered to kill

By employee type (the split that defines Outlay's land-and-expand): - **Engineers / data scientists** → model (b). Raw **API tokens + coding agents**; consumption-based; the **deepest attribution gap** (account-level bills, no team/ ticket metadata). **This is where Outlay lands** — it's the spend that's biggest per head, most volatile, and least visible. [Menlo; Vantage] - **Knowledge workers** → model (a). **Seats**; predictable but **wasteful** (over-buying) and now

drifting to metered AI add-ons. Outlay's value here is **budget/showback + waste reduction** — the **expand** motion across the rest of the company.

The shift that creates the opening: the industry is moving **per-seat** → **consumption/credits** because AI can't be seat-priced [Flexera]. Each new metered tool/provider adds another opaque, account-level bill with no native team/workload attribution → finance can't map credits to work → budget surprises. **Outlay sits across all of them and produces the attributed, forecastable, governable view that no single vendor's billing console gives.** As consumption pricing spreads from eng to every function, the attribution gap — and Outlay's TAM — spreads with it.

5. US TAM model (Outlay analysis — assumptions labeled)

The research provided strong building blocks (LLM API ≈ \$8.4B mid-2025 global; 88% increasing budgets; FinOps \$83B+ sample; capex \$400–450B) but **no clean published "US governable AI-spend" number**, so the pool below is **modeled bottoms-up with explicit assumptions** and sanity-checked top-down. Treat as ranges, not point estimates.

5a. Spend-under-management pool (US, 2026) — what Outlay can govern

Layer	US 2026 (modeled)	Basis / assumption
LLM provider API usage	~\$10–15B	Global enterprise API ≈ \$8.4B mid-2025, annualizing/growing ~2×/yr → ~\$20–30B global 2026; US ≈ ~50% [Menlo]
Employee AI SaaS seats (Copilot, ChatGPT-Ent, Claude, Cursor, Gemini, Glean...)	~\$8–15B	~20–35M US paid AI seats × ~\$300–450/yr; penetration rising fast off 88% budget-increase intent [Wharton]
Company-paid cloud AI inference/training (their own Bedrock/Azure/GCP AI, GPU)	~\$15–35B	Demand-side slice of cloud AI that <i>customer companies</i> pay (excludes hyperscaler/lab capex)
Governable US AI-spend pool	≈ \$35–70B (2026), → ~\$100B+ by 2027–28	Sum; growing ~50–100%/yr [Deloitte ramp; FinOps 98%]

Top-down sanity check: global AI software/services spend estimates run ~\$150–300B in 2026 depending on scope; US ≈ 40–50% → ~\$60–150B all-in, of which the **cleanly governable** SaaS+API+company-cloud slice is the ~\$35–70B above. Consistent order of magnitude. (*medium confidence — wide range by design.*)

5b. Outlay serviceable revenue TAM

FinOps/cost-governance tools monetize as either a **% of spend under management (~1–3%)** or a platform subscription; savings-share models (Outlay's "% of realized savings") tend to net out in a similar **~2–5% of governed spend** band.

	Assumption	US revenue TAM
Total addressable	2–5% take on \$35–70B pool	~\$1.0–3.0B
Serviceable (near-term ICP)	high-fit segment ≈ 25–35% of pool, 2–4% take	~\$0.3–0.9B

	Assumption	US revenue TAM
Obtainable (3-yr beachhead)	low-single-% of serviceable	~\$15–60M ARR is a credible early ceiling for the wedge

5c. Core vs incremental TAM (the roadmap question)

- **Core product TAM — attribution + forecasting + budget governance** (the must-have to capture the base). This is what the FinOps-for-AI buyer needs *first* and what 98%-now-manage-AI-spend are shopping for. **≈ 60–70% of the revenue TAM (~\$0.7–2.1B)**. It's the price of entry and the durable wedge.
- **Incremental-feature TAM — enforcement/guardrails, chargeback/showback, procurement & commitment optimization, anomaly/waste reduction, multi-cloud FinOps-for-AI**. Each is a wallet expander on the same accounts. **≈ 30–40% on top (~\$0.3–0.9B)**, and rising as consumption pricing spreads to every function. Procurement optimization in particular grows as commitments/credits proliferate.
- **Implication: the core must nail attribution + forecast + budget** across the (b) consumption layer first — that's where the unaddressed pain and the biggest TAM slice are. Enforcement/chargeback are expansion, not the wedge.

5d. Fastest-growing, highest-fit segments (where to point GTM)

1. **Eng-heavy growth/late-stage software & AI-native companies (250–5k emp)** — biggest consumption gap, fastest spend growth, can't-yet-build-it. **Primary beachhead**.
2. **AI-forward financial services & insurance** — large bills, mandated cost governance, multi-provider.
3. **Regulated / public sector** — governance is *required* (your Maryland motion); slower cycle, higher ACV, defensible once SOC 2/StateRAMP land.
4. **Digital-native scale-ups in support-heavy verticals** (usage-metered support AI).

6. Customer personas

1. **"Platform Eng Lead at a Series C AI-native SaaS co"** (300 emp, eng-heavy). Buys model (b): central Anthropic/OpenAI/Bedrock orgs + Cursor/Copilot seats. Owns the AI budget himself; finance asks why it 3×'d. **Buys Outlay** to attribute tokens to team/feature, forecast the quarter, and cap runaway agents. *Highest fit*.
2. **"VP Finance / FinOps Lead at a 2,000-person digital enterprise."** Owns the cross-tool budget; AI is now the #1 line he can't explain. Multiple providers, no unified view. **Buys Outlay** for showback/chargeback and budget governance. *High fit — the economic buyer*.
3. **"Head of Eng at a PE-backed scale-up under margin pressure."** PE wants efficiency; AI spend is growing and unattributed. **Buys Outlay** to prove cost-per-output and cut waste. *High fit*.
4. **"CIO/CISO at a regulated enterprise or gov agency"** (your Maryland persona). Needs governance + audit + budget control, metadata-only for compliance. **Buys Outlay** as the AI cost-governance layer. *High ACV, slow cycle*.
5. **"IT/Procurement Lead at a 5k-person traditional enterprise."** Mostly seats today (Copilot, ChatGPT-Ent), lots of shadow AI. **Buys Outlay** for seat-waste + shadow-AI discovery + budget. *Medium fit, large expand*.
6. **(Low fit) "Founder at a 20-person pre-seed."** One ChatGPT Team seat set; bill too small. Not now.

7. Trends expanding the TAM (next 2-3 years)

- **Consumption/credits pricing spreads from eng to every function** → the attribution gap (and Outlay's pool) widens beyond engineering. [Flexera]
 - **Agentic / token-usage explosion** — agents loop and consume far more tokens than chat; inference already two-thirds of compute and rising → recurring spend balloons. [Menlo; Deloitte]
 - **AI-cost governance is institutionalizing as a FinOps sub-discipline** — 98% now manage AI spend, it's the #1 priority and #1 skill gap → a budget line and an owner now exist to sell to. [FinOps 2026]
 - **Multi-provider is the steady state** (Anthropic/OpenAI/Google/Bedrock) → no single billing console will ever be the system of record; a neutral cross-provider layer is needed. [Menlo]
 - **Shadow AI + surprise bills** drive security/finance to demand visibility → governance mandate. [Flexera]
-

8. Caveats, weak spots, and open questions

Caveats / time-sensitivity. - Market-share and spend figures move fast (Menlo mid-year → year-end already shifted Anthropic 32% → 40%); treat point numbers as quarters-fresh. - Several primary pages (Deloitte, FinOps, Wharton, Menlo, Flexera) 403'd to direct fetch; facts were verified via search-index reproduction + independent corroboration (still high-confidence on the quotes, but we couldn't deep-read the full PDFs). - The TAM pool is **modeled**, not a published figure — the research gave building blocks, not a clean US number. Ranges are intentional.

Claims that FAILED verification — do NOT cite these (they sound useful but didn't hold up): - "\$3.5B → \$8.4B doubled in exactly six months" (the doubling-timeframe framing; the \$8.4B level itself is fine). - "53.4% of orgs don't understand the full scope of their AI spend." - "41% of enterprises waste >15% of AI spend; only 7.5% embed FinOps in AI projects." (*refuted 0-3*) - "74% already report positive GenAI ROI." - "Hybrid pricing is the fastest-growing model; +21% median growth." (*refuted 0-3*)

Open questions worth a follow-up run / primary research. 1. **Buy-vs-credits mechanics by employee type, quantified** — public data describes the shift but not the % split of companies on seats vs central-credits vs reimbursement vs shadow, by size. This is a prime **customer-interview** question (and a survey we could own). 2.

Realized AI-spend waste % — the clean "X% is wasted" stat got refuted; we need our own measured number (the dogfood/pilot data is how we earn it). 3. **Willingness-to-pay / take-rate** for AI-cost governance specifically (cloud FinOps benchmarks are the proxy; AI-native pricing is unproven). 4. **Where the budget actually sits** (Eng vs Finance vs IT) by segment — determines who we sell to first.

9. Implications for Outlay

- **Lead with the engineering/consumption wedge.** The biggest, fastest-growing, least-visible spend is eng on **API tokens + coding agents** — and it's exactly what metadata-only/BYOK/read-only fits. Win the platform/eng + FinOps buyer there first.
- **The core product is non-negotiable and is most of the TAM:** attribution → forecast → budget governance across multi-provider consumption spend. Nail that before enforcement/chargeback.
- **Qualify on the \$3 variables;** the threshold (>\$250–500k AI spend, consumption-heavy, multi-tool, no owner, eng-driven) is your ICP filter.
- **The regulated/gov motion** (Maryland) is a real, high-ACV *flavor* of persona #4 — but the **volume beachhead** is eng-heavy growth-stage software (#1–#3). Run both: gov for lighthouse ACV, software for velocity.

- **Earn the proof the research couldn't give us** — measured coverage + waste numbers from real pilots are both the product validation *and* the marketing stat that the refuted claims show is missing.