

US market analysis — AI-spend attribution & governance (Outlay)

Question: Who buys AI in the US, which variables predict they get real value from Outlay (metadata-only, BYOK, read-only attribution + forecasting + budget governance of AI spend), how do they buy AI compute today, and what's the defensible TAM — core vs incremental?

Provenance. Built from a deep-research run (fan-out web search → fetch → **3-vote adversarial verification**). 19 claims survived verification; 5 were refuted and excluded (listed at the end so we don't repeat them). Macro facts below are **cited**; the taxonomy, fit-scoring, purchasing breakdown, and TAM are **Outlay's analysis on top of those facts**, with assumptions labeled. Scope: **US**, demand-side AI spend (what companies pay to use AI — SaaS + API + their own cloud AI), which is the pool Outlay can govern. (Supply-side data-center capex is sized separately and is **not** Outlay's addressable pool — see TAM caveat.)

0. Executive summary

1. **The category is real and inflecting now.** AI-spend management went from a niche to near- universal in 24 months — **98% of organizations now manage AI spend as part of FinOps, up from 31% in 2024**, the fastest cost-discipline adoption the FinOps survey has ever recorded, and it's the **#1 FinOps priority** for the next 12 months. [FinOps Foundation 2026]
 2. **The pain is structural, not cyclical.** AI bills scale by **tokens/compute, not seats**, so per-seat pricing can't price them; spend is **consumption-based, opaque, account-level, and spread across many providers with no unified view** — the exact attribution gap Outlay closes. [Flexera; Datadog; Vantage]
 3. **The bill is large, growing, and finance-visible. 88% of large US enterprises plan to increase GenAI budgets in the next 12 months; 62% by ≥10%.** Enterprise LLM **API** spend alone was **~\$8.4B** by mid-2025 (≈2× in ~6 months), and the field is **multi-vendor** (Anthropic ~32–40%, OpenAI ~25–27%, Google ~20%, Llama ~9%). [Wharton/GBK; Menlo Ventures]
 4. **The wedge is the engineering/coding-agent segment.** Code generation is AI's first breakout use case (**~\$4B, ~28% of enterprise LLM deployments**, Claude ~42% developer share), it runs on **raw API tokens + coding agents** (the worst attribution gap), and **inference/production has overtaken training** — i.e., recurring usage cost, not one-time. [Menlo Ventures]
 5. **TAM (Outlay model).** US governable AI-spend pool **≈ \$35–70B in 2026**, growing ~50–100%/yr. At a FinOps-style **2–5% software take-rate**, that's a **~\$1–3B US revenue TAM**, of which the **core product** (attribution + forecast + budget governance) is the entry wedge and **incremental features** (enforcement, chargeback, procurement optimization) roughly double the wallet over time.
-

1. Market context (the facts the strategy rests on)

AI-cost governance has become a board-level FinOps discipline — fast. - 98% of organizations now manage AI spend as part of FinOps, **up from 31% two years ago** (2024); the survey calls it the fastest technology-cost adoption in its history. [FinOps Foundation, *State of FinOps 2026*, n=1,192 orgs representing \$83B+ tech spend — data.finops.org] (*high confidence*) - **Managing AI spend is the #1 FinOps priority** for the next 12 months (cited by 32.7% of teams — a plurality), and **AI cost management is the #1 skillset** teams want to build (58%). [FinOps

2026] (high) - FinOps scope is expanding beyond cloud: **90% manage (or plan to manage) SaaS, 64% licensing, 57% private cloud, 48% data center.** [FinOps 2026] (high) — i.e., the buyer Outlay sells to already exists and is actively widening their remit to exactly this.

The bill is big, growing, multi-vendor, and consumption-shaped. - 88% of US enterprises (1,000+ employees, >\$50M revenue) plan to increase GenAI budgets in the next 12 months; 62% expect ≥10% increases.

[Wharton/GBK Collective, 2025 AI Adoption Report, ~800 US enterprise decision-makers — ai.wharton.upenn.edu] (high) - Enterprise **LLM API spend ≈ \$8.4B by mid-2025**, roughly double ~6 months prior; basis is a Menlo survey of 150 technical leaders. [Menlo Ventures, 2025 Mid-Year LLM Market Update] (medium — the "\$8.4B" figure is corroborated; the specific "doubled in exactly 6 months" framing failed verification) - **Multi-vendor field:** Anthropic ~32% (year-end ~40%), OpenAI ~25% (down from ~50% in 2023), Google ~20%, Meta Llama ~9%. [Menlo] (high on the numbers; nuance: top-3 ≈ 88% concentration and only ~11% of teams switched providers in the past year — vendor diversity is real at the market level and growing within large enterprises, but isn't yet universal per-company multi-homing) - **Inference (production) has overtaken training** — 74% of startups and 49% of large enterprises say most compute is now inference — i.e., **recurring usage cost**, not one-time training. [Menlo] (high) - **Code generation is the breakout use case** (~\$4B, ~28% of enterprise LLM deployments; Claude ~42% developer share vs OpenAI ~21%). [Menlo] (high)

Why per-seat pricing broke (the structural opening). - AI features scale by tokens, compute-minutes, and model complexity — not by user — so per-seat pricing structurally can't price them; this is the core driver of the **seat → consumption** shift (~65% of vendors now use hybrid seat+usage pricing). [Flexera; corroborated by Bain] (high) - **Consumption data is opaque and aggregated at the account level**, and AI usage is **spread across internal platforms, SaaS tools, and vendor APIs with no out-of-the-box unified view** — making attribution to workloads/teams hard. Provider billing "typically breaks down costs only to the model, project, or workspace level," and "team/cost-center metadata doesn't exist natively in AI billing data." [Flexera; Datadog; Vantage] (high) — **this is Outlay's problem statement, verbatim. - Finance often can't map a credit back to a workload**; credit expirations/resets/rollover "silently create overages or strand value," producing budget surprises and uncontrolled use. A reported **78% of IT leaders saw surprise AI charges** in the past year. [Flexera] (high) - **Seat overbuying already wastes money** (~51–53% of SaaS licenses go unused; large enterprises waste ~\$21M/yr), and AI makes the meter worse. [Flexera; CFO Dive] (high)

Supply-side context (NOT the addressable pool). Global **AI data-center capex is forecast at \$400–450B in 2026** (\$250–300B of it chips), rising to **~\$1T by 2028**; inference ≈ two-thirds of all compute in 2026. [Deloitte, TMT Predictions 2026] (high). This is the hyperscaler/AI-lab **build**; Outlay governs the **consuming company's** AI bill, so capex is the backdrop, not the TAM (see §4 caveat).

2. Customer taxonomy — the variables, and what each implies for fit

Categorize US AI buyers on these axes. For each, the **"hot" end** (higher Outlay value) is bolded.

Variable	Buckets (low-fit → high-fit)	Why it predicts value
AI spend model	seat-SaaS only → mixed → heavy API/ consumption + credits	Attribution gap lives in token/usage spend, not flat seats [Flexera]
Spend magnitude & growth	flat/small → large & growing ≥10%/yr	88% increasing budgets; bigger bill = bigger problem [Wharton]
# providers / tools		Fragmentation = no native rollup [Menlo; Datadog]

Variable	Buckets (low-fit → high-fit)	Why it predicts value
	single tool → many tools, multi-provider, no unified view	
Cost-owner gap	one owner, in control → no owner / finance blind	"can't map credit → workload" [Flexera]
AI maturity / current state	laggard → experimenter → scaled adopter (production inference) → AI-native	Inference = recurring cost; experiments are too small to govern [Menlo]
Employee mix	non-technical, light AI → engineering/data-heavy + AI-native knowledge work	Eng on coding agents = highest token spend [Menlo]
Size (employees)	<50 → 50–250 → 250–5k → 5k+	Big enough to have a real bill + departments to attribute; not so big they've already built it
Stage / type	pre-seed → growth/late-stage startup, PE-backed, established corp digitizing → public mega-cap	Budget pressure + finance scrutiny + can't-yet-build-it-themselves
Vertical	low-AI (e.g. traditional mfg) → software/tech, financial services, digital-native, prof. services, gov/regulated	Where AI spend concentrates and where governance is mandated
Key tools	one chat seat → Cursor/Copilot + Claude/OpenAI/Bedrock APIs + Copilot/ChatGPT-Ent seats	Coding agents + provider APIs are the consumption core

AI usage by employee function (where the spend actually is): - Engineering / data-ML — highest \$/head. Coding agents (Cursor, Copilot, Claude Code) + raw provider API tokens (OpenAI/Anthropic/Bedrock/Azure). Consumption-based, volatile, hardest to attribute. *This is the core wedge.* [Menlo: code gen \$4B/28%] - **Customer support / service** — high and rising (deflection bots, agent assist); +16pp perceived impact vs 2024. Often usage-metered. [Wharton] - **Sales & marketing** — broad seat adoption (Copilot, ChatGPT, content tools); mostly seats today. - **Legal, Procurement, Finance, HR, Ops** — fastest-rising *perceived* impact (Legal **+24pp**, Procurement **+15pp**) — AI spend is **spreading across departments**, multiplying who needs attribution. [Wharton] Mostly seats now, drifting to usage. - Cross-cutting: **82% of enterprise leaders use GenAI weekly, 46% daily** — adoption is broad-based and recurring (though that's leaders, not yet the whole workforce). [Wharton]

3. The critical value-prop-fit variables (ranked by predictive power)

If you could ask only a few questions to qualify a prospect, ask these — in order:

1. **"How much of your AI spend is consumption/usage-based (API tokens, credits, metered) vs flat seats — and is it growing?"** (*highest signal.*) Heavy + growing consumption spend = the attribution gap is real and worsening. [Flexera; Menlo]
2. **"How many AI tools/providers are in use, and can anyone produce one unified view of the spend and who/what drove it?"** Many tools + "no, nobody can" = textbook Outlay. [Datadog; Vantage]
3. **"Who owns the AI budget, and have you been surprised by an AI bill?"** No clear owner + surprise charges (78% have) = active pain, budget authority to buy. [Flexera]
4. **"What's your AI maturity — experiments, or production workloads running every day?"** Production inference = recurring, governable spend; experiments are too small. [Menlo]
5. **"What share of your headcount is engineering/data, and do they use coding agents / call model APIs directly?"** Eng-heavy = the highest-spend, lowest-visibility function. [Menlo]

6. **Spend magnitude / company size** — gates whether the ROI clears your price (below a threshold the bill is too small to bother governing).

The threshold where the value prop becomes real (Outlay's qualifying line):

A company with **>\$250–500k/yr of AI spend, at least a meaningful share consumption-based, 2+ AI tools/providers, no single owner producing an attributed view, and an engineering/AI-native function** driving usage. Below that, the bill is too small or too simple to need a governance layer (seats with one owner ≈ low fit).

High-fit ICP: growth/late-stage software & digital-native companies (250–5k employees), AI-native or "scaled adopter," engineering-heavy, multi-provider API + coding-agent spend, finance starting to ask "what is this costing per team/project?" Plus **AI-forward financial services and regulated/gov** where governance is *mandated* (your existing Maryland/gov motion is the regulated flavor of this). **Low-fit:** pre-seed/tiny (no bill), AI-laggards (no spend), single-seat-SaaS shops with one owner (little attribution pain — though budget/waste value still exists).

4. How companies buy AI compute today (the central question)

Four models coexist; most companies run **several at once**, which is itself the problem.

Model	Who owns budget	Pricing	Who/what it covers	Why they choose it	Outlay value
(a) Seat / enterprise subscriptions (ChatGPT Enterprise, Copilot, Claude for Work, Glean)	IT / Procurement	per-seat, predictable	knowledge workers (sales, mktg, ops, legal, support)	control, security/SSO, simple budgeting, compliance	budget/showback + waste (overbought seats, ~50% unused)
(b) Central API/ compute + credits/ budgets (OpenAI/ Anthropic/Bedrock/ Azure orgs; cloud GPU)	Engineering / Platform / sometimes Finance	usage/ consumption , volatile	engineers, data/ML, production features, coding agents	speed + developer experience, can't seat-price tokens, scale	core wedge — attribute tokens → team/ticket/ workload, forecast, cap
(c) Reimbursed individual subscriptions	Finance (expensed)	per-seat, scattered	early/SMB, individual power users	speed, low friction, pre-policy	near-zero visibility; consolidation/ showback opportunity
(d) Shadow AI on personal accounts/ cards	nobody (ungoverned)	mixed	anyone, pre-policy	velocity, avoiding procurement	security/compliance + surprise-charge risk; the thing governance teams are chartered to kill

By employee type (the split that defines Outlay's land-and-expand): - **Engineers / data scientists** → model (b). Raw **API tokens + coding agents**; consumption-based; the **deepest attribution gap** (account-level bills, no team/ ticket metadata). **This is where Outlay lands** — it's the spend that's biggest per head, most volatile, and least visible. [Menlo; Vantage] - **Knowledge workers** → model (a). **Seats**; predictable but **wasteful** (over-buying) and now

drifting to metered AI add-ons. Outlay's value here is **budget/showback + waste reduction** — the **expand** motion across the rest of the company.

The shift that creates the opening: the industry is moving **per-seat** → **consumption/credits** because AI can't be seat-priced [Flexera]. Each new metered tool/provider adds another opaque, account-level bill with no native team/workload attribution → finance can't map credits to work → budget surprises. **Outlay sits across all of them and produces the attributed, forecastable, governable view that no single vendor's billing console gives.** As consumption pricing spreads from eng to every function, the attribution gap — and Outlay's TAM — spreads with it.

5. US TAM model (Outlay analysis — assumptions labeled)

The research provided strong building blocks (LLM API ≈ \$8.4B mid-2025 global; 88% increasing budgets; FinOps \$83B+ sample; capex \$400–450B) but **no clean published "US governable AI-spend" number**, so the pool below is **modeled bottoms-up with explicit assumptions** and sanity-checked top-down. Treat as ranges, not point estimates.

5a. Spend-under-management pool (US, 2026) — what Outlay can govern

Layer	US 2026 (modeled)	Basis / assumption
LLM provider API usage	~\$10–15B	Global enterprise API ≈ \$8.4B mid-2025, annualizing/growing ~2×/yr → ~\$20–30B global 2026; US ≈ ~50% [Menlo]
Employee AI SaaS seats (Copilot, ChatGPT-Ent, Claude, Cursor, Gemini, Glean...)	~\$8–15B	~20–35M US paid AI seats × ~\$300–450/yr; penetration rising fast off 88% budget-increase intent [Wharton]
Company-paid cloud AI inference/training (their own Bedrock/Azure/GCP AI, GPU)	~\$15–35B	Demand-side slice of cloud AI that <i>customer companies</i> pay (excludes hyperscaler/lab capex)
Governable US AI-spend pool	≈ \$35–70B (2026), → ~\$100B+ by 2027–28	Sum; growing ~50–100%/yr [Deloitte ramp; FinOps 98%]

Top-down sanity check: global AI software/services spend estimates run ~\$150–300B in 2026 depending on scope; US ≈ 40–50% → ~\$60–150B all-in, of which the **cleanly governable** SaaS+API+company-cloud slice is the ~\$35–70B above. Consistent order of magnitude. (*medium confidence — wide range by design.*)

5b. Outlay serviceable revenue TAM

FinOps/cost-governance tools monetize as either a **% of spend under management (~1–3%)** or a platform subscription; savings-share models (Outlay's "% of realized savings") tend to net out in a similar **~2–5% of governed spend** band.

	Assumption	US revenue TAM
Total addressable	2–5% take on \$35–70B pool	~\$1.0–3.0B
Serviceable (near-term ICP)	high-fit segment ≈ 25–35% of pool, 2–4% take	~\$0.3–0.9B

	Assumption	US revenue TAM
Obtainable (3-yr beachhead)	low-single-% of serviceable	~\$15–60M ARR is a credible early ceiling for the wedge

5c. Core vs incremental TAM (the roadmap question)

- **Core product TAM — attribution + forecasting + budget governance** (the must-have to capture the base). This is what the FinOps-for-AI buyer needs *first* and what 98%-now-manage-AI-spend are shopping for. **≈ 60–70% of the revenue TAM (~\$0.7–2.1B)**. It's the price of entry and the durable wedge.
- **Incremental-feature TAM — enforcement/guardrails, chargeback/showback, procurement & commitment optimization, anomaly/waste reduction, multi-cloud FinOps-for-AI**. Each is a wallet expander on the same accounts. **≈ 30–40% on top (~\$0.3–0.9B)**, and rising as consumption pricing spreads to every function. Procurement optimization in particular grows as commitments/credits proliferate.
- **Implication: the core must nail attribution + forecast + budget** across the (b) consumption layer first — that's where the unaddressed pain and the biggest TAM slice are. Enforcement/chargeback are expansion, not the wedge.

5d. Fastest-growing, highest-fit segments (where to point GTM)

1. **Eng-heavy growth/late-stage software & AI-native companies (250–5k emp)** — biggest consumption gap, fastest spend growth, can't-yet-build-it. **Primary beachhead**.
2. **AI-forward financial services & insurance** — large bills, mandated cost governance, multi-provider.
3. **Regulated / public sector** — governance is *required* (your Maryland motion); slower cycle, higher ACV, defensible once SOC 2/StateRAMP land.
4. **Digital-native scale-ups in support-heavy verticals** (usage-metered support AI).

6. Customer personas

1. **"Platform Eng Lead at a Series C AI-native SaaS co"** (300 emp, eng-heavy). Buys model (b): central Anthropic/OpenAI/Bedrock orgs + Cursor/Copilot seats. Owns the AI budget himself; finance asks why it 3×'d. **Buys Outlay** to attribute tokens to team/feature, forecast the quarter, and cap runaway agents. *Highest fit*.
2. **"VP Finance / FinOps Lead at a 2,000-person digital enterprise."** Owns the cross-tool budget; AI is now the #1 line he can't explain. Multiple providers, no unified view. **Buys Outlay** for showback/chargeback and budget governance. *High fit — the economic buyer*.
3. **"Head of Eng at a PE-backed scale-up under margin pressure."** PE wants efficiency; AI spend is growing and unattributed. **Buys Outlay** to prove cost-per-output and cut waste. *High fit*.
4. **"CIO/CISO at a regulated enterprise or gov agency"** (your Maryland persona). Needs governance + audit + budget control, metadata-only for compliance. **Buys Outlay** as the AI cost-governance layer. *High ACV, slow cycle*.
5. **"IT/Procurement Lead at a 5k-person traditional enterprise."** Mostly seats today (Copilot, ChatGPT-Ent), lots of shadow AI. **Buys Outlay** for seat-waste + shadow-AI discovery + budget. *Medium fit, large expand*.
6. **(Low fit) "Founder at a 20-person pre-seed."** One ChatGPT Team seat set; bill too small. Not now.

7. Trends expanding the TAM (next 2-3 years)

- **Consumption/credits pricing spreads from eng to every function** → the attribution gap (and Outlay's pool) widens beyond engineering. [Flexera]
 - **Agentic / token-usage explosion** — agents loop and consume far more tokens than chat; inference already two-thirds of compute and rising → recurring spend balloons. [Menlo; Deloitte]
 - **AI-cost governance is institutionalizing as a FinOps sub-discipline** — 98% now manage AI spend, it's the #1 priority and #1 skill gap → a budget line and an owner now exist to sell to. [FinOps 2026]
 - **Multi-provider is the steady state** (Anthropic/OpenAI/Google/Bedrock) → no single billing console will ever be the system of record; a neutral cross-provider layer is needed. [Menlo]
 - **Shadow AI + surprise bills** drive security/finance to demand visibility → governance mandate. [Flexera]
-

8. Caveats, weak spots, and open questions

Caveats / time-sensitivity. - Market-share and spend figures move fast (Menlo mid-year → year-end already shifted Anthropic 32% → 40%); treat point numbers as quarters-fresh. - Several primary pages (Deloitte, FinOps, Wharton, Menlo, Flexera) 403'd to direct fetch; facts were verified via search-index reproduction + independent corroboration (still high-confidence on the quotes, but we couldn't deep-read the full PDFs). - The TAM pool is **modeled**, not a published figure — the research gave building blocks, not a clean US number. Ranges are intentional.

Claims that FAILED verification — do NOT cite these (they sound useful but didn't hold up): - "\$3.5B → \$8.4B doubled in exactly six months" (the doubling-timeframe framing; the \$8.4B level itself is fine). - "53.4% of orgs don't understand the full scope of their AI spend." - "41% of enterprises waste >15% of AI spend; only 7.5% embed FinOps in AI projects." (*refuted 0-3*) - "74% already report positive GenAI ROI." - "Hybrid pricing is the fastest-growing model; +21% median growth." (*refuted 0-3*)

Open questions worth a follow-up run / primary research. 1. **Buy-vs-credits mechanics by employee type, quantified** — public data describes the shift but not the % split of companies on seats vs central-credits vs reimbursement vs shadow, by size. This is a prime **customer-interview** question (and a survey we could own). 2.

Realized AI-spend waste % — the clean "X% is wasted" stat got refuted; we need our own measured number (the dogfood/pilot data is how we earn it). 3. **Willingness-to-pay / take-rate** for AI-cost governance specifically (cloud FinOps benchmarks are the proxy; AI-native pricing is unproven). 4. **Where the budget actually sits** (Eng vs Finance vs IT) by segment — determines who we sell to first.

9. Implications for Outlay

- **Lead with the engineering/consumption wedge.** The biggest, fastest-growing, least-visible spend is eng on **API tokens + coding agents** — and it's exactly what metadata-only/BYOK/read-only fits. Win the platform/eng + FinOps buyer there first.
- **The core product is non-negotiable and is most of the TAM:** attribution → forecast → budget governance across multi-provider consumption spend. Nail that before enforcement/chargeback.
- **Qualify on the \$3 variables;** the threshold (>\$250–500k AI spend, consumption-heavy, multi-tool, no owner, eng-driven) is your ICP filter.
- **The regulated/gov motion** (Maryland) is a real, high-ACV *flavor* of persona #4 — but the **volume beachhead** is eng-heavy growth-stage software (#1–#3). Run both: gov for lighthouse ACV, software for velocity.

- **Earn the proof the research couldn't give us** — measured coverage + waste numbers from real pilots are both the product validation *and* the marketing stat that the refuted claims show is missing.

10. Commitment & procurement optimization — the highest-\$ incremental lever

Why it matters. AI vendors price like cloud: an **on-demand rack rate**, plus **committed-spend discounts** (commit \$X/yr → ~10–40% off) and **provisioned throughput / reserved capacity** (Azure OpenAI PTUs, Bedrock provisioned, reserved GPUs — a fixed price for a dedicated lane). Committing lowers the **unit price**, but the **effective** cost depends on utilization: there's a **break-even** above which committing wins and below which on-demand is cheaper (you avoid paying for idle/forfeited capacity). Self-hosting open models on reserved GPUs is the extreme — cheapest per token only at huge, steady scale. In cloud FinOps this commitment/Reserved-Instance optimization is the **single highest-dollar lever**; reported AI-infra savings run **30–50% on GPU RIs/SPs and 30–60% in two quarters** across GPU/LLM/inference. [nOps; Opslyft; ProsperOps]

Why companies get it wrong (the opening). To land above break-even you must **forecast usage and attribute it** well enough to size the commit. Most can't, so they (a) overpay on-demand (never negotiated), (b) over-commit and forfeit, or (c) under-commit and eat overage rates. Each cloud's PTU/CUD sales motion individually looks reasonable but **together quietly locks in over-commitment** — the discipline is to decide **at the portfolio level after ≥30 days of real traffic data**, which is exactly the attributed + forecasted history Outlay produces.

Feature sketch (incremental, built on the core data): - **Commitment recommender** — from the attributed usage curve: "X% of your Anthropic/OpenAI/Bedrock spend is steady → commit/provision that baseline; leave the spiky remainder on-demand," with a modeled effective-savings-rate and break-even. - **Commitment pacing** — are you on track to *use* your committed spend, or forfeit it? (a direct extension of the program/budget pacing already shipped). - **PTU vs on-demand split by workload** — which teams/features are steady enough to move to provisioned throughput vs leave metered. - **Cheaper-equivalent levers, advisory (not in-path)** — flag batch-API / caching / cheaper-model opportunities by workload (the routing/caching that gateways do live; Outlay *recommends* it from the read-only data without sitting in the request path). - **Renewal / negotiation pack** — the attributed spend + forecast as leverage at contract time.

This is the clearest path from **core (attribution + forecast + budget)** to a **high-ROI incremental TAM**, mapping onto the proven cloud-FinOps commitment-optimization business — but at the model-API layer the cloud optimizers don't cover.

11. Competitive landscape — who optimizes AI-compute spend today

Three layers; **no one owns the full stack**, but the core-attribution layer is getting crowded and Flexera is consolidating.

Layer	Players	What they do	Gap vs Outlay
Cloud commitment optimizers	ProsperOps (acq. by Flexera, Jan 2026), Zesty, nOps	Autonomously ladder RIs / Savings Plans / CUDs to maximize discount on cloud	Optimize infra you run , not LLM API commitments / model-layer PTUs ; no work-attribution . Most enterprises are on managed APIs they don't cover.

Layer	Players	What they do	Gap vs Outlay
		GPU instances (~\$5B claimed savings; 30–50% on GPU)	
AI cost visibility / FinOps (nearest core competitors)	Vantage, CloudZero, Finout , Datadog, Flexera FinOps-for-AI , usage.ai, Amnic	Connect OpenAI/Anthropic/cloud billing; per-request + model-level attribution; map spend to team/project/customer; anomaly detection; some forecasting	Attribution mostly at team/project/customer level, not the work-item (ticket/feature) join ; cloud-first with AI bolt-ons; none lead with metadata-only/BYOK/read-only or forecast-against-delivered-work
LLM gateways / routers	LiteLLM, Portkey, OpenRouter , Cloudflare AI Gateway, TrueFoundry , Helicone, Martian	In-path request optimization: caching (20–40% savings), cheaper-model routing, fallback, per-key/per-team budget caps	Sit in the data path (opposite of Outlay's posture); platform-eng buyer ; do request-level cuts + budget enforcement, not commitment/portfolio optimization or financial governance
Provider / cloud-native dashboards	OpenAI/Anthropic usage, Azure/AWS cost tools	Per-vendor usage + billing	Single-vendor, aggregated, no cross-provider view, no commitment intelligence
SaaS-spend / procurement	Vendr, Tropic, Zylo, Productiv, Ramp	Negotiate/manage SaaS contracts + seats	Touch AI seat licenses, not token/compute commitment optimization

Reads for Outlay: - Differentiation must be sharp — the *visibility/attribution* core (Vantage/CloudZero/Finout) is contested. Outlay's defensible edges: (1) **work-item attribution** via the tracker join (token → ticket → engineer → feature), (2) **forecasting back-tested on the customer's own delivered work**, (3) **metadata-only / BYOK / read-only** (the compliance/regulated wedge none of them lead with), (4) **budget governance + pacing** as a first-class workflow, not just dashboards. - **The commitment-optimization layer (§10) is open at the model-API level** — but **Flexera (now owning ProsperOps + FinOps-for-AI) is the most likely to stitch it together**. Move while it's a gap; the work-attribution data is the moat that makes recommendations better than an infra-only optimizer. - **Gateways are a posture choice, not just a competitor** — they enforce in-path; Outlay governs out-of-path (read-only). For regulated/compliance buyers, out-of-path is the selling point; for hard-enforcement needs, a gateway integration (or the shipped enforcement endpoint) is the bridge.

Sources (§10–11): ProsperOps · ProsperOps/Flexera, Zesty, nOps · nOps LLM cost tools · Vantage · CloudZero · Finout — AI cost observability · Holori top-10 AI cost tools · LLM gateways · Portkey gateway guide · Opslyft AI cost optimization 2026