Outlay.ai

Put AI compute on a budget.

Prepared for Maryland Health Connection

Krethikram Gowrisankar · Founder, Outlay
June 26, 2026 · outlay-ai.com

ABOUT

KG

San Francisco, CA

Krethikram Gowrisankar

Founder, Outlay · ex-Amazon (Senior PM, Tech)

- Senior PM (Tech) at Amazon — led Smart TDR, an autonomous docking system scaled to ~50% of the middle-mile network in 12 months: 1M+ automated cycles, zero safety incidents.
- Delivered $40M+ annual savings and owned $50M+ annual investment / $100M+ NPV programs — cost governance and capital allocation at scale.
- Repeat 0→1 builder across hardware, software, and operations; filed a provisional patent and cold-pitched Amazon leadership into a full-time PM role.
- B.S. Computer Science, UW–Madison — built Outlay to bring that same financial rigor to AI/LLM spend.

krethikram@gmail.com · outlay-ai.com

OUTLAY · FOR MARYLAND HEALTH CONNECTION

AI spend, mapped to the work that drove it — and held to budget.

A read-only, metadata-only governance layer for the AI behind Maryland Health Connection.

Attribute · Forecast · Govern

THE MOMENT

Maryland is operationalizing its 2025 AI Enablement Strategy.

Right after “is it safe?” comes the question every CFO and CIO asks:

> **“What is the AI costing us — per program — and is it on budget?”**

Today that lives in invoices and spreadsheets, after the money is spent. Outlay answers it continuously, attributed to the work.

WHAT OUTLAY DOES

Three things, from data you already have.

- **Attribute** — every AI dollar mapped to a team, work type, ticket, and engineer.
- **Forecast** — what planned work will cost, back-tested on your delivered work.
- **Govern** — budgets per program, real-time pacing, on-/off-track alerts.

Read-only links — no agents, nothing in the request path.

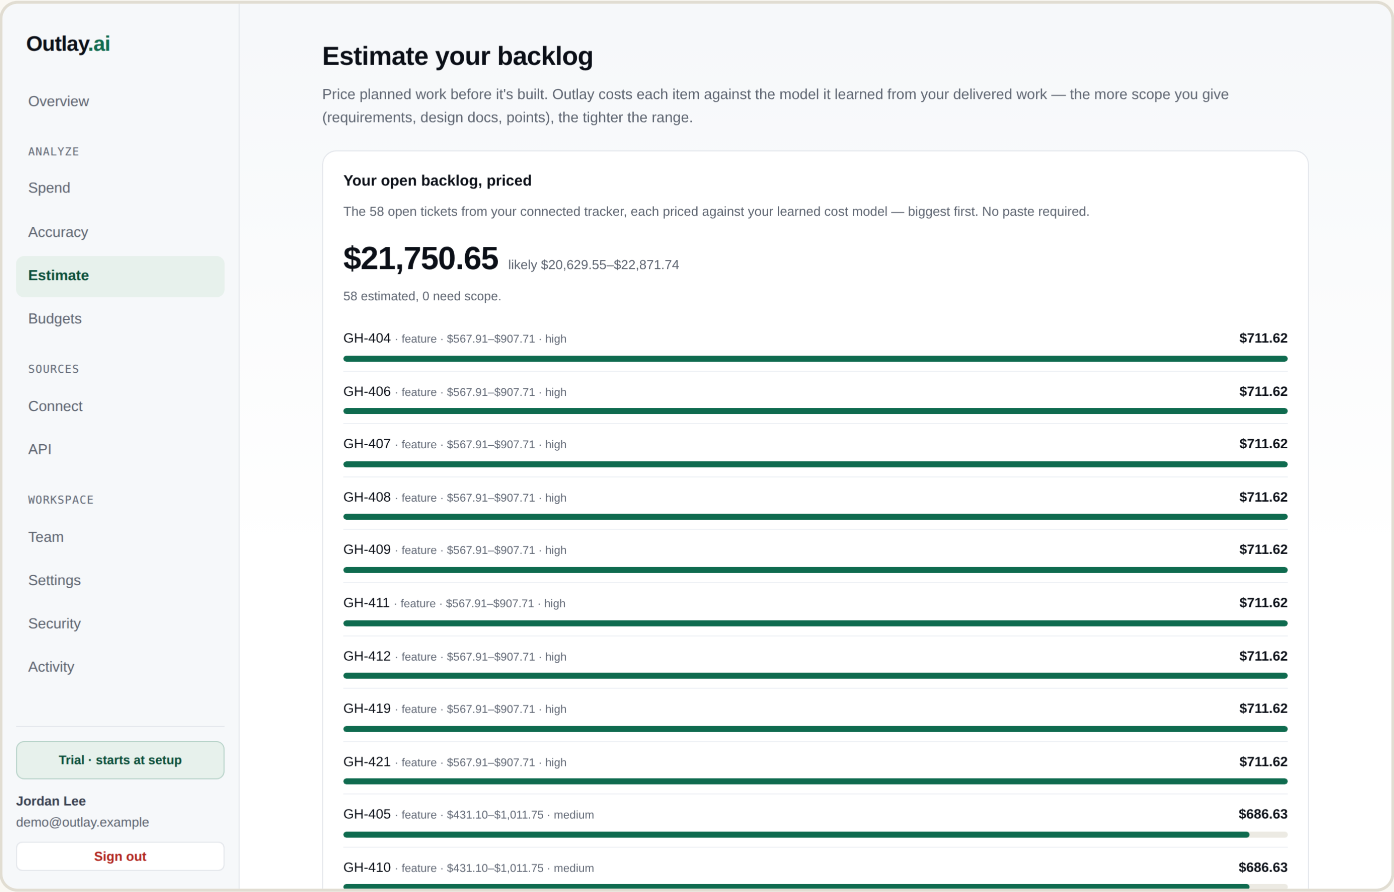

WHY THIS IS SAFE FOR A HEALTH EXCHANGE

We never see PHI, PII, or FTI — by design.

- **Metadata-only** — token counts, ticket IDs, $ — never prompts or outputs.
- **Bring-your-own-key** — keys stay in your environment.
- **Read-only** — never in a data path.
- Out of **MARS-E / 1075 / HIPAA** data scope — what stops most vendors.

How Outlay connects — metadata only, nothing in the data path

Read-only links observe usage and cost. Your data never leaves your environment.

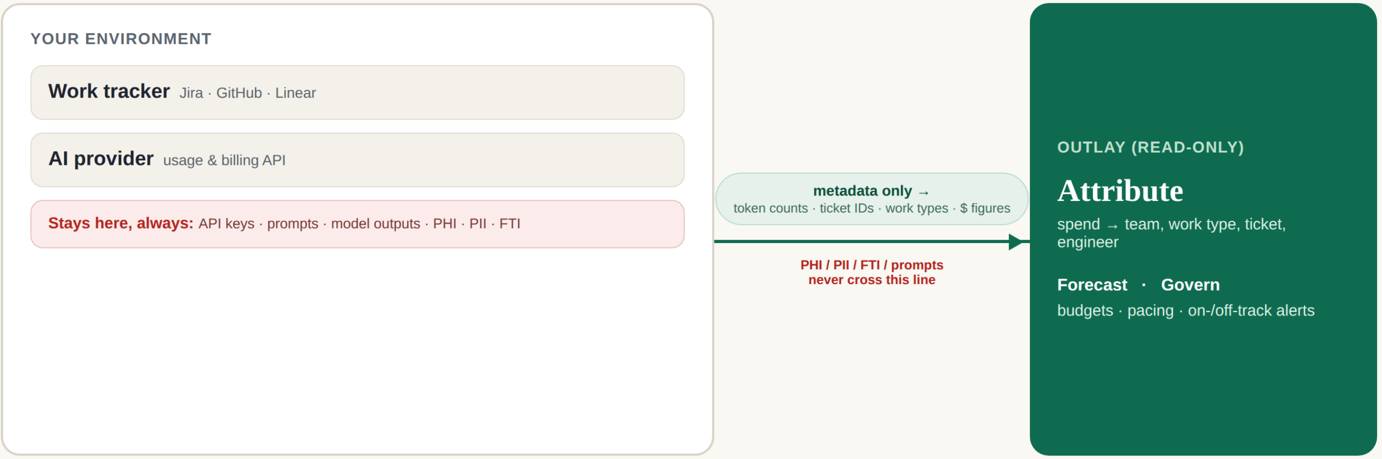

LIVE

Let’s look at the product.

On sample data shaped like a real engineering program — the same screens you’d see on day one of a pilot.

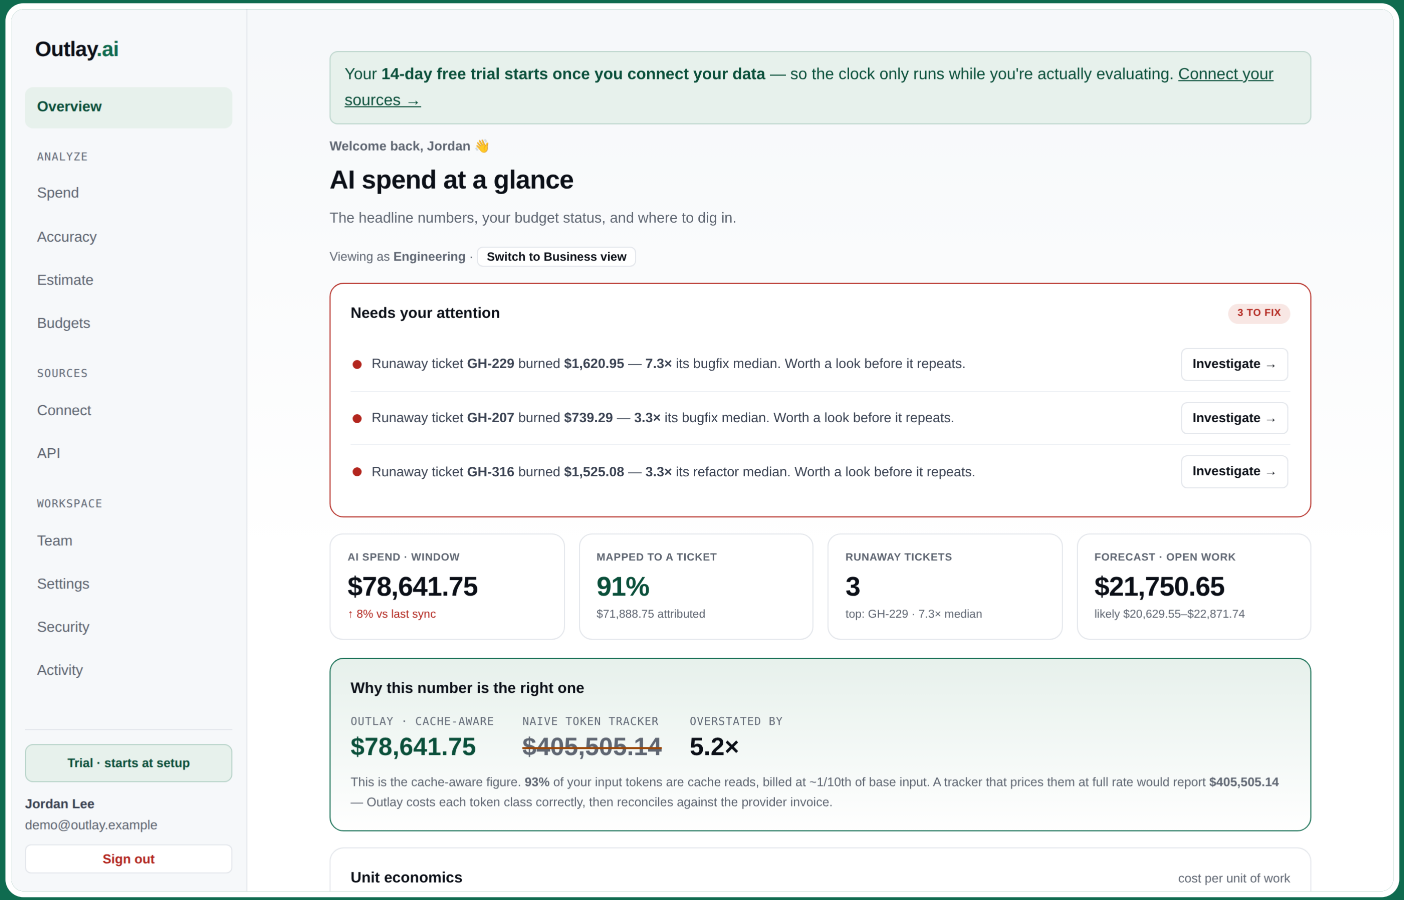

DEMO · ATTRIBUTION

Every dollar, on the roadmap.

- By team & cost center, work type, ticket, engineer.
- Ticket coverage shown honestly — and how to lift it.
- FOCUS-aligned CSV export for finance.

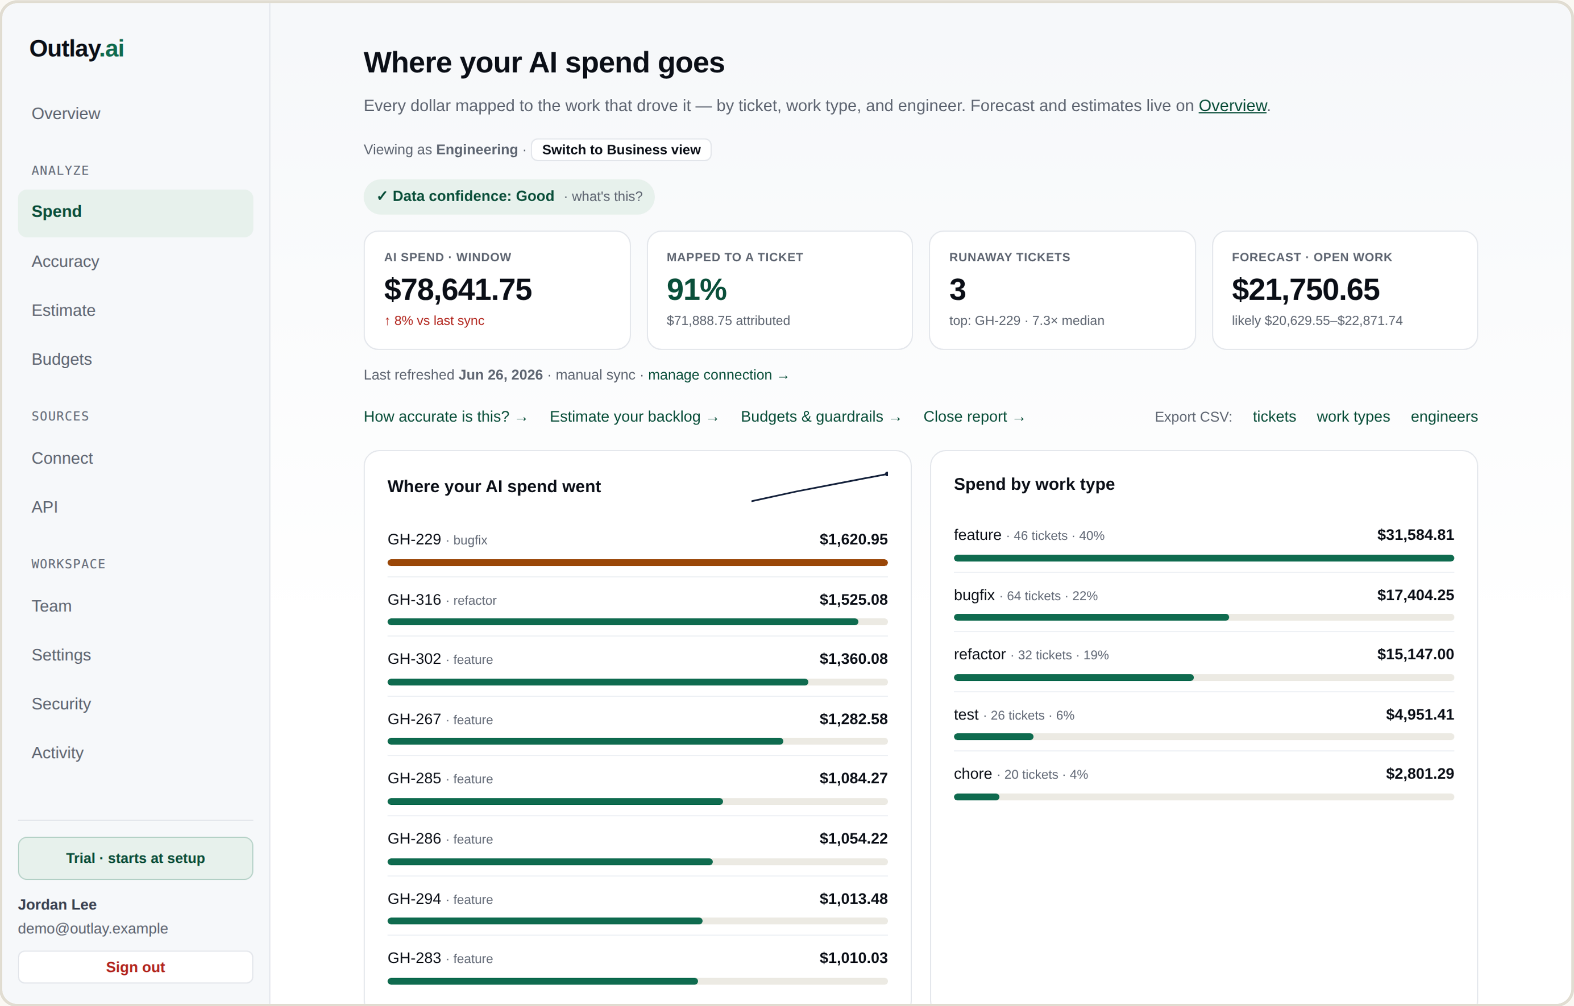

DEMO · ACCURACY (THE HONESTY LAYER)

We don’t ask you to trust a benchmark.

- Back-tested on your own closed tickets (leave-one-out).
- Median error, within-P90, sample size — never hidden.
- Per-work-type accuracy with over/under bias.

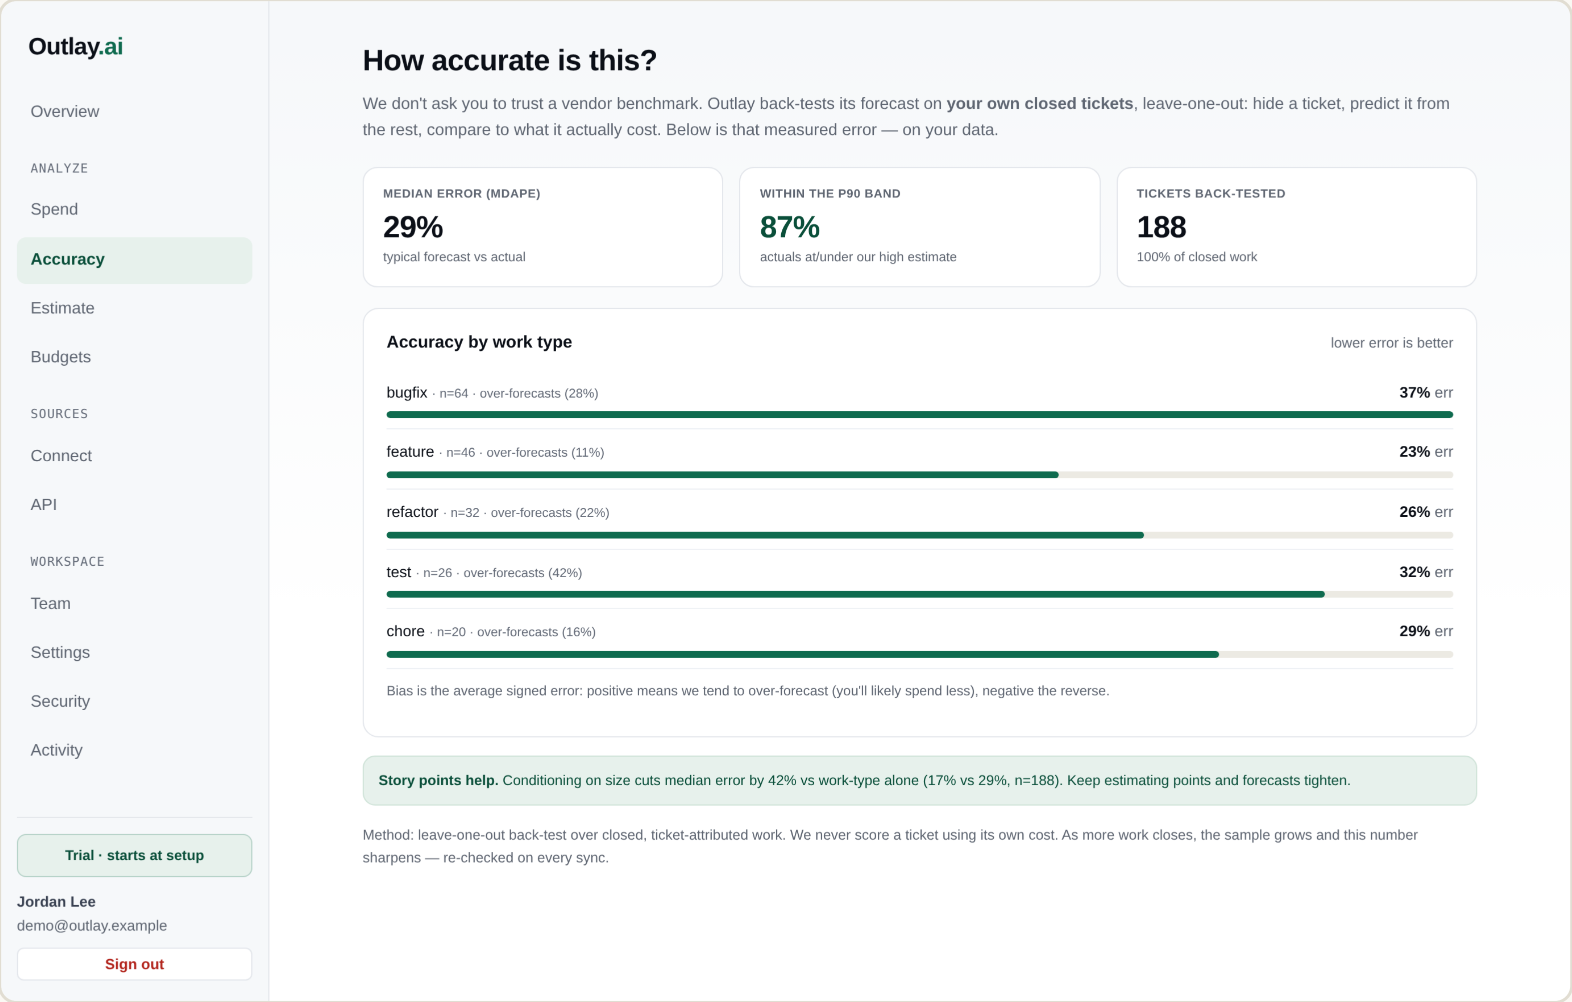

DEMO · GOVERNANCE

Put AI compute on a budget — per program.

- Program budgets across teams, projects, work types.
- Real-time pacing → the projected breach date.
- On-/off-track earned-value rating; alerts to your SIEM.

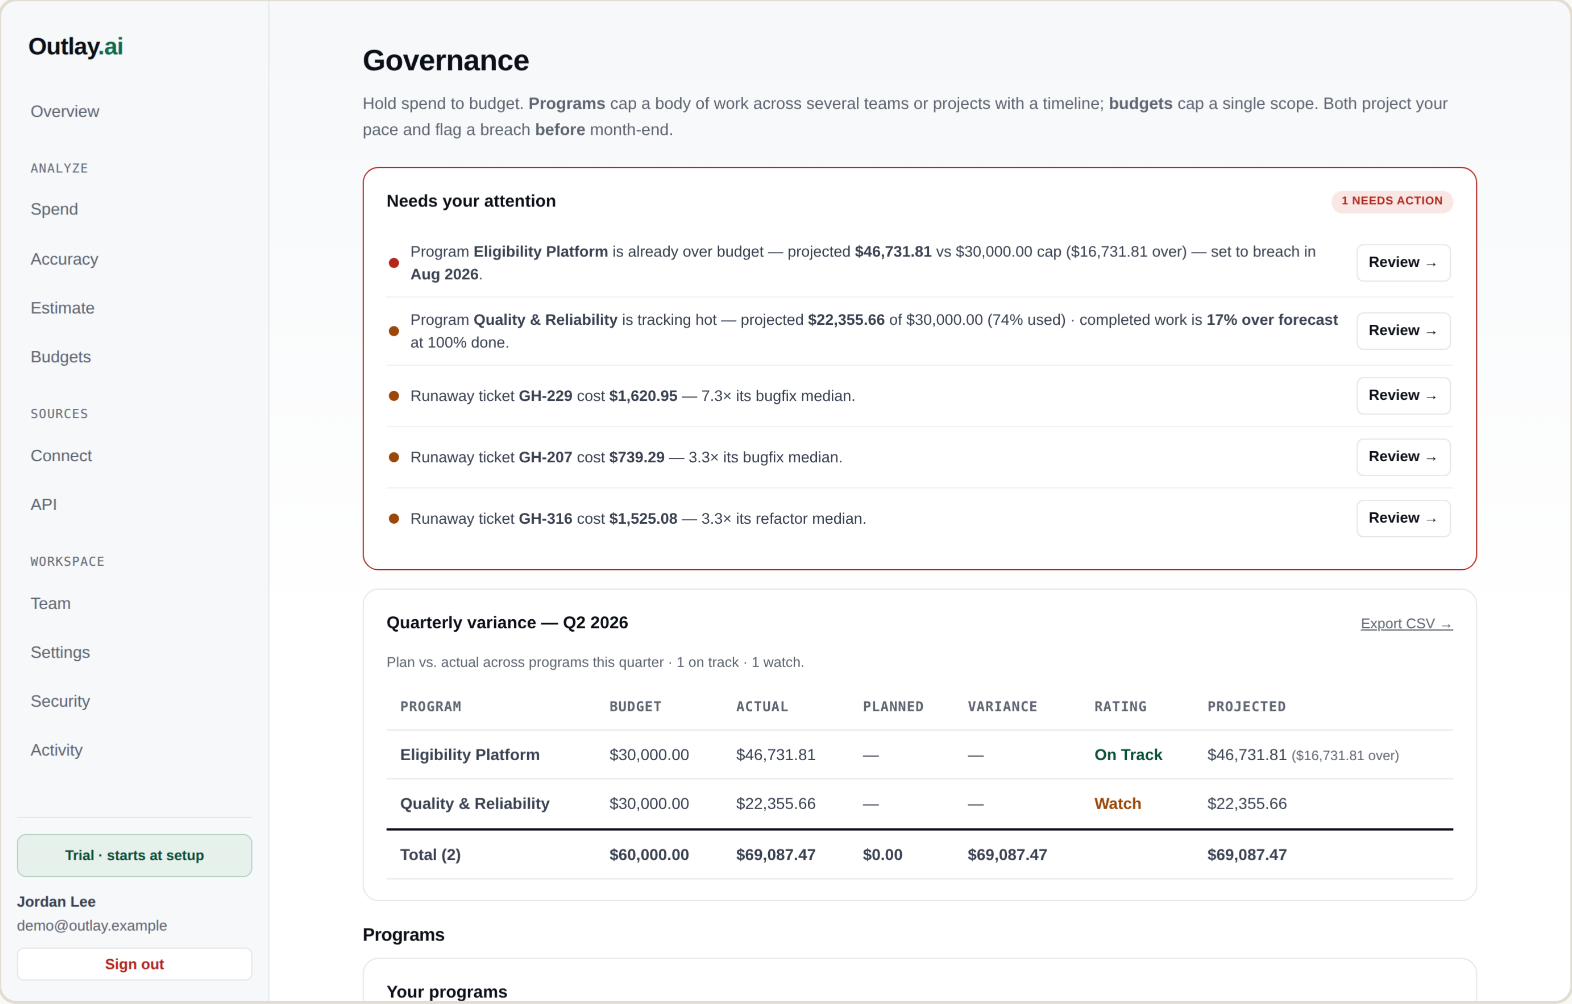

SECURITY & COMPLIANCE POSTURE

Straight talk on where we are.

IN PLACE TODAY

- SSO (OIDC) + SCIM provisioning
- Phishing-resistant MFA (passkeys)
- RBAC · audit log → SIEM export
- Retention + self-serve erasure
- Encryption at rest + KMS hook
- IR plan · WCAG 2.1 AA · CI security scan

ON THE ROADMAP (HONEST)

- SOC 2 Type II — in progress
- Annual penetration test — scheduled
- StateRAMP / GovRAMP — roadmap
- FedRAMP cloud + FIPS crypto — deal-gated
- Metadata-only keeps most of these out of scope.

FIT FOR MHBE

Two ways this earns its place now.

THE AI-GOVERNANCE LAYER

- Be the cost-attribution + budget-guardrail layer as you stand up the 2025 AI strategy.
- Governance-first — matched to how your program is built.
- Day-one value, no instrumentation: read-only.

VENDOR & ENHANCEMENT SPEND

- As Deloitte / J29 work adopts AI coding agents, see cost per project.
- The earned-value view a public agency needs to justify IT spend.
- Attributes the $12.7M “MHC enhancements” to the tickets behind it.

THE ASK

A two-week, read-only pilot — zero data-handling risk.

- Scoped to your **development / vendor AI usage**, not any production PHI system.
- Metadata-only and read-only — nothing for the security review to clear on data flow.
- You walk away with a real number: **attribution coverage** + a **measured forecast accuracy**.

Free during the pilot · we share our SOC 2 timeline and NIST 800-53 control mapping up front.

WHY NOW

Govern the AI spend before it scales.

The cheapest time to put AI on a budget is at the start of the curve — exactly where Maryland is.

Krethikram Gowrisankar · krethikram@gmail.com · outlay-ai.com