

AI spend, mapped to the work that drove it — **and held to budget.**

A read-only, metadata-only governance layer for the AI behind Maryland Health Connection.

Attribute · Forecast · Govern

THE MOMENT

Maryland is operationalizing its 2025 AI Enablement Strategy.

Right after “is it safe?” comes the question every CFO and CIO asks:

“What is the AI costing us — per program — and is it on budget?”

Today that lives in invoices and spreadsheets, after the money is spent. Outlay answers it continuously, attributed to the work.

WHAT OUTLAY DOES

Three things, from data you already have.

- **Attribute** — every AI dollar mapped to a team, work type, ticket, and engineer.
- **Forecast** — what planned work will cost, back-tested on your delivered work.
- **Govern** — budgets per program, real-time pacing, on-/off-track alerts.

Read-only links — no agents, nothing in the request path.

Outlay.ai

Overview

ANALYZE

Spend

Accuracy

Estimate

Budgets

SOURCES

Connect

API

WORKSPACE

Team

Settings

Security

Activity

Trial · starts at setup

Jordan Lee
demo@outlay.example

Sign out

Estimate your backlog

Price planned work before it's built. Outlay costs each item against the model it learned from your delivered work — the more scope you give (requirements, design docs, points), the tighter the range.

Your open backlog, priced

The 58 open tickets from your connected tracker, each priced against your learned cost model — biggest first. No paste required.

\$21,750.65

likely \$20,629.55–\$22,871.74

58 estimated, 0 need scope.

GH-404 · feature · \$567.91–\$907.71 · high	\$711.62
GH-406 · feature · \$567.91–\$907.71 · high	\$711.62
GH-407 · feature · \$567.91–\$907.71 · high	\$711.62
GH-408 · feature · \$567.91–\$907.71 · high	\$711.62
GH-409 · feature · \$567.91–\$907.71 · high	\$711.62
GH-411 · feature · \$567.91–\$907.71 · high	\$711.62
GH-412 · feature · \$567.91–\$907.71 · high	\$711.62
GH-419 · feature · \$567.91–\$907.71 · high	\$711.62
GH-421 · feature · \$567.91–\$907.71 · high	\$711.62
GH-405 · feature · \$431.10–\$1,011.75 · medium	\$686.63
GH-410 · feature · \$431.10–\$1,011.75 · medium	\$686.63

WHY THIS IS SAFE FOR A HEALTH EXCHANGE

We never see PHI, PII, or FTI — by design.

- **Metadata-only** — token counts, ticket IDs, \$ — never prompts or outputs.
- **Bring-your-own-key** — keys stay in your environment.
- **Read-only** — never in a data path.
- Out of **MARS-E / 1075 / HIPAA** data scope — what stops most vendors.

How Outlay connects — metadata only, nothing in the data path

Read-only links observe usage and cost. Your data never leaves your environment.

YOUR ENVIRONMENT

Work tracker Jira · GitHub · Linear

AI provider usage & billing API

Stays here, always: API keys · prompts · model outputs · PHI · PII · FTI

metadata only →
token counts · ticket IDs · work types · \$ figures

**PHI / PII / FTI / prompts
never cross this line**

OUTLAY (READ-ONLY)

Attribute

spend → team, work type, ticket, engineer

Forecast · Govern

budgets · pacing · on-/off-track alerts

LIVE

Let’s look at the product.

On sample data shaped like a real engineering program — the same screens you’d see on day one of a pilot.

Outlay.ai

Overview

ANALYZE

Spend

Accuracy

Estimate

Budgets

SOURCES

Connect

API

WORKSPACE

Team

Settings

Security

Activity

Trial · starts at setup

Jordan Lee

demo@outlay.example

Sign out

Your 14-day free trial starts once you connect your data — so the clock only runs while you’re actually evaluating. [Connect your sources](#) →.

Welcome back, Jordan 🙌

AI spend at a glance

The headline numbers, your budget status, and where to dig in.

Viewing as Engineering · [Switch to Business view](#)

Needs your attention

Runaway ticket GH-229 burned \$1,620.95 — 7.3× its bugfix median. Worth a look before it repeats.

Investigate →

Runaway ticket GH-207 burned \$739.29 — 3.3× its bugfix median. Worth a look before it repeats.

Investigate →

Runaway ticket GH-316 burned \$1,525.08 — 3.3× its refactor median. Worth a look before it repeats.

Investigate →

AI SPEND · WINDOW

\$78,641.75

↑ 8% vs last sync

MAPPED TO A TICKET

91%

\$71,888.75 attributed

RUNAWAY TICKETS

3

top: GH-229 · 7.3× median

FORECAST · OPEN WORK

\$21,750.65

likely \$20,629.55–\$22,871.74

Why this number is the right one

OUTLAY · CACHE-AWARE

NAIVE TOKEN TRACKER

OVERSTATED BY

\$78,641.75

~~\$405,505.14~~

5.2×


This is the cache-aware figure. 93% of your input tokens are cache reads, billed at ~1/10th of base input. A tracker that prices them at full rate would report \$405,505.14 — Outlay costs each token class correctly, then reconciles against the provider invoice.

Unit economics

cost per unit of work

Every dollar, on the roadmap.

- By team & cost center, work type, ticket, engineer.
- Ticket coverage shown honestly — and how to lift it.
- FOCUS-aligned CSV export for finance.



We don't ask you to trust a benchmark.

- Back-tested on your own closed tickets (leave-one-out).
- Median error, within-P90, sample size — never hidden.
- Per-work-type accuracy with over/under bias.

Outlay.ai

Overview

ANALYZE

Spend

Accuracy

Estimate

Budgets

SOURCES

Connect

API

WORKSPACE

Team

Settings

Security

Activity

Trial · starts at setup

Jordan Lee
demo@outlay.example

Sign out

How accurate is this?

We don't ask you to trust a vendor benchmark. Outlay back-tests its forecast on **your own closed tickets**, leave-one-out: hide a ticket, predict it from the rest, compare to what it actually cost. Below is that measured error — on your data.

MEDIAN ERROR (MDAPE)

29%

typical forecast vs actual

WITHIN THE P90 BAND

87%

actuals at/under our high estimate

TICKETS BACK-TESTED

188

100% of closed work

Accuracy by work type

lower error is better

bugfix · n=64 · over-forecasts (28%)

37% err

feature · n=46 · over-forecasts (11%)

23% err

refactor · n=32 · over-forecasts (22%)

26% err

test · n=26 · over-forecasts (42%)

32% err

chore · n=20 · over-forecasts (16%)

29% err

Bias is the average signed error: positive means we tend to over-forecast (you'll likely spend less), negative the reverse.

Story points help. Conditioning on size cuts median error by 42% vs work-type alone (17% vs 29%, n=188). Keep estimating points and forecasts tighten.

Method: leave-one-out back-test over closed, ticket-attributed work. We never score a ticket using its own cost. As more work closes, the sample grows and this number sharpens — re-checked on every sync.

Put AI compute on a budget — per program.

- Program budgets across teams, projects, work types.
- Real-time pacing → the projected breach date.
- On-/off-track earned-value rating; alerts to your SIEM.

Outlay.ai

Overview

ANALYZE

Spend

Accuracy

Estimate

Budgets

SOURCES

Connect

API

WORKSPACE

Team

Settings

Security

Activity

Trial · starts at setup

Jordan Lee
demo@outlay.example

Sign out

Governance

Hold spend to budget. **Programs** cap a body of work across several teams or projects with a timeline; **budgets** cap a single scope. Both project your pace and flag a breach **before** month-end.

Needs your attention1 NEEDS ACTION

●

Program **Eligibility Platform** is already over budget — projected **\$46,731.81** vs \$30,000.00 cap (\$16,731.81 over) — set to breach in **Aug 2026**.

Review --

●

Program **Quality & Reliability** is tracking hot — projected **\$22,355.66** of \$30,000.00 (74% used) · completed work is **17% over forecast** at 100% done.

Review --

●

Runaway ticket **GH-229** cost **\$1,620.95** — 7.3× its bugfix median.

Review --

●

Runaway ticket **GH-207** cost **\$739.29** — 3.3× its bugfix median.

Review --

●

Runaway ticket **GH-316** cost **\$1,525.08** — 3.3× its refactor median.

Review --

Quarterly variance — Q2 2026Export CSV --

Plan vs. actual across programs this quarter · 1 on track · 1 watch.

PROGRAM	BUDGET	ACTUAL	PLANNED	VARIANCE	RATING	PROJECTED
Eligibility Platform	\$30,000.00	\$46,731.81	—	—	On Track	\$46,731.81 (\$16,731.81 over)
Quality & Reliability	\$30,000.00	\$22,355.66	—	—	Watch	\$22,355.66
Total (2)	\$60,000.00	\$69,087.47	\$0.00	\$69,087.47		\$69,087.47

Programs

Your programs

Straight talk on where we are.

IN PLACE TODAY

- SSO (OIDC) + SCIM provisioning
- Phishing-resistant MFA (passkeys)
- RBAC · audit log → SIEM export
- Retention + self-serve erasure
- Encryption at rest + KMS hook
- IR plan · WCAG 2.1 AA · CI security scan

ON THE ROADMAP (HONEST)

- SOC 2 Type II — in progress
- Annual penetration test — scheduled
- StateRAMP / GovRAMP — roadmap
- FedRAMP cloud + FIPS crypto — deal-gated
- Metadata-only keeps most of these out of scope.

Two ways this earns its place now.

THE AI-GOVERNANCE LAYER

- Be the cost-attribution + budget-guardrail layer as you stand up the 2025 AI strategy.
- Governance-first — matched to how your program is built.
- Day-one value, no instrumentation: read-only.

VENDOR & ENHANCEMENT SPEND

- As Deloitte / J29 work adopts AI coding agents, see cost per project.
- The earned-value view a public agency needs to justify IT spend.
- Attributes the \$12.7M “MHC enhancements” to the tickets behind it.

THE ASK

A two-week, read-only pilot — zero data-handling risk.

- Scoped to your **development / vendor AI usage**, not any production PHI system.
- Metadata-only and read-only — nothing for the security review to clear on data flow.
- You walk away with a real number: **attribution coverage** + a **measured forecast accuracy**.

Free during the pilot · we share our SOC 2 timeline and NIST 800-53 control mapping up front.

WHY NOW

Govern the AI spend before it scales.

The cheapest time to put AI on a budget is at the start of the curve — exactly where Maryland is.