Outlay — Outbound playbook (eng-consumption ICP)

How to find, target, and open the beachhead from `docs/product-strategy.md`: **eng-heavy software / AI-native companies, 250–5,000 employees, >$250–500k/yr consumption-based AI spend, multi-provider, with no clear owner.** Founder-led outbound; the goal of every touch is a 15-minute diagnostic call, not a demo.

The core insight that makes outbound work here: the pain (AI bill growing 50–100%/yr that no one can attribute) and our sharpest differentiator (the work-item attribution join) peak in the same accounts. We don't need to create pain — we name a bill they're already nervous about and offer the one answer they can't get from a dashboard: what is each team/feature actually costing?

---

1. Targeting — who exactly

Firmographic filters (the list): - **Industry:** software/SaaS, AI-native/ML products, dev-tools, fintech-with-heavy-eng. - **Size:** 250–5,000 employees; **eng headcount ≥ ~80** (proxy for AI-tool spend at scale). - **Stage:** Series B–D or profitable scale-up — has a finance function and a real bill. - **Geo:** US first (matches the market analysis + compliance posture).

Buying-signal filters (rank the list — these predict pain): 1. **Public AI/LLM usage** — ships an AI feature, or eng blog/conference talk about LLM infra (they have token spend). 2. **Multi-provider footprint** — mentions of Anthropic and OpenAI/Bedrock/Azure OpenAI, plus Cursor/Copilot rollouts (the attribution gap is worst here). 3. **Hiring signal** — open reqs for "FinOps," "Platform/Infra eng," or "AI platform" → someone was told to get costs under control. 4. **Funding signal** — raised in last 12 mo → board now asks about burn/efficiency; AI is the fastest-growing line. 5. **Tooling signal** — job posts / stack pages listing Jira/Linear + GitHub (our join lands cleanly).

Who to contact (two-buyer motion — message both, different angle): - **Economic/Finance:** VP Finance, Head of FinOps, CFO (smaller orgs). Angle: attribution + forecast + budget control. - **Technical/Champion:** VP Eng, Head of Platform/Infra, Staff Platform Eng. Angle: stop being the one who can't answer "what did that cost," metadata-only so no security review.

Disqualify before you spend a touch (from the battlecards): mostly self-hosted GPU on cloud (ProsperOps fits better), AI spend still trivial (few $20 seats), or no work tracker / no branch hygiene (our join degrades). Don't burn the account — note and revisit.

---

2. The sequence (4 touches over ~12 days, multi-channel)

Keep emails **<120 words**, one ask, no attachments on touch 1, plain text. Personalize the first line from a real signal (their AI feature, a talk, a job post) — never "I came across your company."

Touch 1 — Email (Day 1) · the named-pain open — to the technical champion

> **Subject:** what did [their AI feature] cost to run last month?
>
> Hi [First] — saw [specific: "your launch of AI-assisted X" / "the talk on your LLM stack"].
>
> Quick question most platform teams can't answer cleanly: **what did that feature cost in model spend last month — and which team drove it?** Provider bills aggregate at the account level, so it usually takes a spreadsheet and a guess.
>
> Outlay joins your AI usage to the work — ticket → engineer → feature — so the answer is one view. **Read-only, metadata-only** (prompts and keys never leave your box, so there's nothing for security to clear).
>
> Worth a 15-min look at your own numbers? — [Name]

Touch 2 — LinkedIn connect + note (Day 3)

> [First] — sent you a note on the AI-cost-attribution problem. We map model spend to the ticket/ feature/engineer (read-only, metadata-only). Happy to share what it surfaces on a stack like yours — no pitch, just the view finance keeps asking eng for.

Touch 3 — Email reply-in-thread (Day 6) · the proof/forecast angle

> Following up with the part finance cares about: once spend is attributed, we **forecast the backlog's AI cost and back-test the error on your own delivered work** — so next quarter's number has an error bar, not a finger in the air. And you can put a **program budget** on a body of work and get a projected-breach date before month-end.
>
> If [VP Finance/their boss] is feeling the "why is the AI bill up again" conversation, this is the 15 minutes that ends it. Open to it this week?

Touch 4 — Email (Day 12) · the break-up + commitment hook

> Last note, [First]. If you're past ~$250k/yr in model spend, the other lever is **how you pay** — committed-spend discounts vs. on-demand. We size that from your run-rate (and flag forfeit risk), which is usually 15–30% no one's claimed yet.

> If now's not the time, reply "later" and I'll check back next quarter. If it is, here's 15 min: [link].

Finance-buyer variant of Touch 1 (when leading with the economic buyer):

> **Subject:** the AI line nobody can explain
>
> Hi [First] — the fastest-growing line in most software P&Ls right now is model/LLM spend, and it's the one finance can't break down by team or initiative (provider bills are account-level totals).
>
> Outlay attributes it to the work — team, feature, engineer — forecasts the next quarter on your own history, and lets you set program budgets that flag a breach before it happens. Read-only, metadata-only.
>
> 15 minutes on your own numbers? — [Name]

---

3. The 15-minute diagnostic call (not a demo)

Goal: qualify + create the "I need this" moment with their data, then scope a pilot. 1. **(2 min) Frame:** "I'll keep this to 15. Goal is to see if the attribution gap is real for you — if not, I'll say so." 2. **(5 min) Diagnose — ask, don't pitch:** - "How many model providers/AI tools are you on now? Who can tell you the per-team split today?" - "Last time finance asked 'why is the AI bill up,' how long did the answer take?" - "Any budget owner on AI spend, or is it diffuse?" - "Roughly what's the monthly run-rate, and how fast is it growing?" 3. **(5 min) Show the one wedge view** (only if qualified): spend → ticket/feature attribution on sample-shaped data, then the program-pacing breach projection. Keep it to the two things they can't get elsewhere. 4. **(3 min) Close to a pilot:** "The honest way to prove this is your own data. A pilot is read-only, metadata-only, ~[X] days to a measured attribution-coverage number and a forecast back-test. Who else needs to be in the room?"

Pilot definition (the real ask): read-only connect (tracker + provider usage), N days, deliver (a) attribution coverage %, (b) a forecast back-test error on their completed work, (c) one program budget with a projected-breach date. Success = a number they didn't have before. (Gov/regulated runs the same motion at higher ACV with the compliance posture as the lead — parallel, not the wedge.)

---

4. Qualification scorecard (BANT-ish, AI-spend-specific)

Signal	Strong (pursue)	Weak (deprioritize)
Monthly model spend	>$20–40k/mo, growing	<$5k/mo, flat
Providers/tools	3+ (multi-vendor gap)	1, single console suffices

Signal	Strong (pursue)	Weak (deprioritize)
Owner	"no one really owns it"	dedicated FinOps already tooled
Tracker hygiene	Jira/Linear + branch discipline	no tracker / no branch links
Trigger	recent raise, board cost pressure, AI launch	none
Hosting	managed APIs (Anthropic/OpenAI/ Bedrock/Azure)	self-hosted GPU on cloud (→ ProsperOps)

3+ "strong" → prioritize. Mostly "weak" → nurture list, revisit next quarter.

---

5. Metrics & cadence (founder-led)

- **Volume:** 10–15 well-researched accounts/week beats 200 blasted. Personalization is the moat at this stage.
- **Funnel targets to watch:** reply rate (aim ≥8–12% on a tight list), call-booked rate, pilot-start rate, pilot→paid.
- **Instrument the message:** track which open (named-feature-cost vs. finance-line vs. commitment) books calls; double down on the winner.
- **Always disqualify out loud** — telling a bad-fit prospect "you don't need us yet" is the highest-trust move and comes back as referrals.

Keep every claim in outreach substantiated — coverage % and forecast error are measured on the prospect's own data in the pilot, never asserted as a benchmark up front.