

Outlay — pricing proposal (concrete packaging)

Turns the analysis in `docs/pricing-research.md` into a concrete, sellable structure: named tiers, price bands, what's in each, the savings-share add-on terms, and the regulated/gov premium. **These are proposed anchors** — final numbers get set with the first 5-10 customers (per `TODO.md`), but a concrete starting point beats negotiating from a blank page.

The decided model (from the research): a **hybrid** — a tiered **platform fee on AI spend under management (SUM)**, sold as predictable flat annual bands, plus an **opt-in savings-share add-on** tied to enforcement/commitment savings. **No per-seat.**

1. The value metric: AI spend under management (SUM)

We bill on the **annualized AI/LLM spend Outlay attributes and governs** — the thing whose value scales with the customer's pain. It's measurable from day one (we compute it), aligns price to value, and is the FinOps-category norm. The **% of SUM declines as spend grows** (like Finout/CloudZero), because the marginal value flattens and large accounts won't pay a flat percentage.

Why not the alternatives (recap): **per-seat** punishes the team-wide read-only visibility that makes us sticky and caps value at a few buyers; **pure %-of-savings** pays \$0 until a customer enables enforcement (most won't on day one) and underprices the always-on attribution/forecast/governance.

2. Platform tiers (annual, on AI spend under management)

Tier	AI spend under mgmt (SUM)	Platform fee / yr	Effective % at top of band	Best for
Team	up to \$500k	\$15,000	3.0%	first owner, single-stack, getting attribution + a forecast
Growth	\$500k - \$2M	\$36,000	1.8%	multi-provider, program budgets + governance in active use
Scale	\$2M - \$10M	\$90,000	0.9%	finance-owned, chargeback + commitment optimization
Enterprise	\$10M+	Custom (≈0.5-0.8%)	—	portfolio commitment optimization, SSO/SCIM at scale, premium support

Notes - Bands are **flat** within the range (predictable for finance/gov budgeting), stepping at the thresholds — not a live percentage meter. - Effective % lands in the FinOps band (visibility tools ≈ 1-3% of spend) at small scale and **below** it at large scale, justified by the

deeper value (work-item attribution + forecast + governance + commitment sizing) over pure visibility. - **Annual prepay** assumed; monthly available at +15%. **Multi-year**: 10% off (2yr), 15% (3yr).

What's included at every tier (the always-on core — don't fragment it): work-item attribution, own-data forecast + back-test, program budgets + pacing + variance, the Commitments recommender (advisory), cross-provider ingest, read-only/BYOK/metadata-only posture, audit logging + SIEM export, SSO/OIDC, 2FA/passkeys, RBAC.

What steps up with tier: SCIM provisioning + enforced MFA (Growth+), commitment **pacing & renewal/negotiation pack** (Scale+), portfolio/multi-provider commitment optimization + premium support SLA + custom retention (Enterprise).

3. The savings-share add-on (opt-in)

Keep the "we only win when you win" hook — but as an **add-on**, not the only model.

- **Terms: 20% of measured net savings** Outlay causes — from (a) the enforcement gateway (route-down / caps, measured against a held-out **control arm**) and (b) realized **commitment** savings the recommender sized. Billed quarterly in arrears on **measured** savings only.
 - **Differentiator vs. ProsperOps: no clawback, no early-termination penalty.** ProsperOps claws back unrealized savings-share for up to 12 months — a documented friction. We don't. Cancel the add-on anytime; you keep the savings.
 - **Disarms the skeptic / lands pilots**: "Turn on enforcement and we only bill if a control arm proves we saved you money." It is additive to the platform fee, gated to where savings are real and attributable.
 - **Guardrail**: never bill savings we didn't measure. The control arm / realized-commitment measurement is the integrity of the model (and a moat — owning a measured savings number the market lacks).
-

4. Regulated / government premium

Same platform, priced for the cost of serving a compliance buyer.

- **+25-50% on the platform fee** once the posture is sold as a SKU (the metadata-only / BYOK / read-only architecture that removes most control scope — see [docs/soc2-stateramp-sequencing.md](#)).
- Bundles: signed security questionnaire + Trust Center evidence, configurable retention, dedicated IR contact, and (as they land) **SOC 2 Type II → StateRAMP/GovRAMP** authorization.

- Higher ACV, longer cycle — run as a **parallel lighthouse** motion (Maryland), not the velocity wedge. Gate the premium on real certs; don't charge for compliance you haven't shipped.
-

5. Pilots & discounting

- **Pilot:** free, time-boxed (~2-4 weeks), read-only/metadata-only. Deliverable = a measured attribution-coverage % + a forecast back-test on the customer's own data. Converts to **Team** or **Growth** on a paid annual.
 - **Design-partner discount:** up to **40% off year 1** for the first ~5-10 logos in exchange for a reference + product feedback. Sunset it deliberately — it's customer-acquisition spend, not the price.
 - **Founder guardrails:** discount the first year, not the list price; never discount the savings-share % (it's already aligned); hold the line on the always-on core being unfragmented (don't sell "attribution-only" cheap — it commoditizes the wedge).
-

6. Worked examples

- **Series-B dev-tools co, ~\$1.2M/yr AI spend, multi-provider, no owner → Growth, \$36k/yr** (3.0% effective). Enables enforcement in Q2 → savings-share add-on on measured cuts. Likely a \$36k + variable ACV.
 - **Scale-up, ~\$6M/yr AI spend, finance-owned, wants chargeback + commitment sizing → Scale, \$90k/yr** (1.5% effective). Renewal/negotiation pack drives a committed-spend deal → savings-share on the realized commitment savings.
 - **State agency, ~\$800k/yr, compliance-led → Growth + gov premium** (≈\$45-54k/yr) once SOC 2 / StateRAMP evidence is in hand; pilot first.
-

7. Open decisions for the founder

1. **Anchor %** — are the 3.0% → 0.9% effective rates right for the value, or anchor lower to win the first logos fast and raise on renewal?
2. **Savings-share %** — 20% vs. the 15-35% category range; lower (15%) is a sharper "we only win when you win" wedge, higher captures more where we clearly cause savings.
3. **Core fragmentation** — keep attribution+forecast+governance bundled (recommended) vs. an entry "attribution-only" SKU to lower the entry price (risks commoditizing the wedge).
4. **SUM definition edge cases** — annualize from a trailing window? include seat-based AI (Copilot) in SUM or only consumption? (Recommend: consumption SUM is the metric; seats are context.)

5. **Publish vs. "talk to us"** — publish Team transparently (self-serve trust) and gate Growth+ to a call? The savings calculator (/savings) already primes the spend-based framing.

All figures proposed/illustrative; validate willingness-to-pay with the first cohort before hardening list prices. Keep every savings claim measured, never asserted.