

Outlay — demo runbook (software / AI-native ICP)

The talk-track for the live product demo to the eng-consumption ICP. Pairs with the outbound playbook (`docs/outbound-playbook.md`) and the leave-behind (`docs/software-icp-leave-behind.pdf`) to make a complete sales kit: **outbound → demo → leave-behind**.

Demo thesis: don't tour features. Run one narrative — "AI spend is a black box → here's the box opened → here's the budget under control → here's how to pay less" — and let two moments land: (1) **a dollar mapped to a ticket**, and (2) **a forecast back-tested on their own work**. Everything else is supporting cast. Target 20-25 minutes + discussion.

0. Before the call (5 min setup)

- Load the **enterprise sample** (`python -m outlay.fixtures.gen_demo`) — ~\$79k/quarter across 8 teams, with the \$79k-real-vs-\$406k-naive callout. A believable bill makes the numbers land.
 - Have **two programs** seeded so Governance shows real pacing (one on-track, one off).
 - Hide demo chrome (sample banners, onboarding cards) if screen-sharing.
 - Open tabs in narrative order: **Overview → Spend → Accuracy → Governance → Commitments**.
 - Re-confirm the prospect's stack from discovery (providers, tracker) so you can say "this is where your Jira/Anthropic data would land."
-

1. Frame (1 min) — name the pain, set the rule

"I'll keep this to ~20 minutes. The goal isn't a feature tour — it's to show you the one question your provider bills can't answer, then answer it. If it's not a real gap for you, I'll tell you."

Anchor on their number from discovery: "You mentioned ~\$X/mo across Anthropic and OpenAI with no clean per-team split — that's exactly what this opens up."

2. Overview (2 min) — the black box, quantified

- Show the consolidated spend with the **real-vs-naive callout** (\$79k vs \$406k).
- Talking point: "First thing — we cost it right. Cache-aware, reconciled to the invoice. Most tools over-count cached tokens 5-10×; if the starting number's wrong, everything downstream is."

- Don't linger — this establishes credibility, then move to the aha.

3. Spend → attribution (5 min) — THE aha #1

- Drill from total → team → **work type** → **a single ticket** with its dollar cost and a **fidelity tier** on it.
- The line to say slowly: "This is the thing no dashboard does. Your provider bill stops at the account. We go token → ticket → engineer → feature. Here's what shipping this feature cost."
- Hit **fidelity**: "Every dollar carries how confident the join is — call, branch, session, team. We'll never hand you a confident per-ticket number that's actually a guess. Finance trusts 'team-level $\pm 15\%$ '; they never forgive a precise wrong number."
- **Pause here**. Let them react. This is the moment the demo is for.

4. Accuracy → forecast back-test (4 min) — THE aha #2 (the trust unlock)

- Show the **leave-one-out back-test**: forecast vs actual on completed work, with the measured error.
- The line: "Anyone can draw a trend line. We forecast your backlog and then prove the error on your own delivered work — leave-one-out. So next quarter's number has an error bar, not a finger in the air. This is measured on your data in the pilot, not a benchmark we're asserting."
- Honesty beat: "Allocating the past is exact; forecasting the future is a range — and we tell you exactly how good it is." (Mirrors the /accuracy page — credibility, not bravado.)

5. Governance → programs (4 min) — from report to system of record

- Show a **program budget** with real-time pacing, the **projected-breach date**, and the **earned-value on-/off-track** rating.
- The line: "Now it's not a report you read — it's a budget you act on. Cap a body of work across teams, and we flag the breach weeks before the invoice, with a date. Here's one on track, here's one that's going to blow its estimate."
- Mention read-only posture in passing: "All of this is read-only and metadata-only — nothing sits in your request path, prompts and keys never leave your box. Nothing for security to clear."

6. Commitments (3 min) — how to pay less (the expansion hook)

- Show the **Commitments** page: spend profile + the committed-spend scenarios with **forfeit risk**, the **provisioned** directional read, and **Export negotiation pack**.
- The line: "Once we know your steady run-rate, we size whether a committed-spend discount beats on-demand — and quantify the forfeit risk so you don't over-commit. Export it and take it to your vendor. This is usually 15-30% nobody's claimed."
- Note the opportunities card: "And we flag caching/batch candidates — advisory, you implement."

7. Close to a pilot (2 min) — the only real ask

"The honest way to prove all of this is your own data. A pilot is **read-only, metadata-only**, ~2 weeks, and you walk away with: your attribution-coverage %, a forecast back-test error on your completed work, and one program budget with a projected-breach date. If the gap isn't real, you'll see that too. Who else needs to be in the room — and can we connect a read-only tracker token this week?"

Leave the one-pager ([software-icp-leave-behind.pdf](#)). Send the recap same day.

Objection handling (the common five)

- **"We already have a cost dashboard (Vantage/CloudZero)."** → "Keep it for cloud and team rollups. We're the layer deeper on the AI line — ticket/feature attribution and a forecast you can defend. Teams run both." (See battlecards §1.)
- **"Our gateway (LiteLLM/Portkey) already caps budgets."** → "In-path caps are great for enforcement. They can't attribute to a feature, forecast the backlog, or size a commitment — and many teams can't put a gateway in the path for compliance. We complement it." (Battlecards §2.)
- **"Is this accurate?"** → Show the back-test again. "That's the whole point — we measure it on your data instead of claiming a number."
- **"Security review will take forever."** → "Metadata-only, BYOK, read-only — no prompts, outputs, or keys leave your environment. That removes most of the review scope. Here's the Trust Center."
- **"We could build this ourselves."** → "You can build a v1 dashboard. The cost is the upkeep — the spend→work join across changing providers, the forecast back-test, the governance workflow. We maintain that so your engineers ship product." (Battlecards §4.)

Demo hygiene

- **Two ahas, max.** Attribution-to-a-ticket and the forecast back-test. If you only land those two, it was a good demo.
- **Don't feature-dump.** Skip Estimate/Connect/Settings unless asked — offer them as "happy to go deeper after."
- **Use their numbers** wherever the UI allows; "your stack here" beats a generic sample.
- **Stay honest.** Every accuracy/savings claim is "measured on your data in the pilot," never asserted. That honesty is the differentiation — don't undercut it by overclaiming live.