

Know what your AI spend is actually buying.

Map every AI dollar to the ticket, feature & engineer · forecast it · hold it to budget.

Outlay.ai

Your model bill is the fastest-growing line in engineering — and the one nobody can break down. Provider invoices aggregate at the account level, so *"what did shipping that feature cost, and which team drove it?"* has no answer. Outlay gives you that answer, forecasts the next quarter, and governs the budget — **read-only**, **metadata-only**.

WHAT IT DOES

- **Attribute** — AI spend mapped to the ticket, feature, engineer & program via the work-tracker join.
- **Forecast** — price the backlog, back-tested on your *own* delivered work (measured error, not a guess).
- **Govern** — per-program budgets, real-time pacing, projected-breach dates, on-/off-track ratings.
- **Commit smart** — size on-demand vs. committed-spend vs. provisioned, with forfeit risk quantified.

HOW IT CONNECTS

- **Trackers** — Jira, Linear, GitHub Issues (read-only).
- **AI usage** — Anthropic, OpenAI, Bedrock, Azure OpenAI, Cursor, Copilot.
- **BYOK** — your keys stay yours; we read usage metadata only.

WHY IT'S DIFFERENT

- **The work-item join** — token → ticket → engineer → feature. Dashboards stop at team/project.
- **Your-data forecast** — leave-one-out error on your history, not a vendor benchmark.
- **Read-only & metadata-only** — prompts, outputs & keys never leave your box. Nothing for security to clear.
- **Commitment sizing** — the model-API layer the cloud optimizers don't cover.

WHO IT'S FOR

- Eng-heavy software / AI-native teams, ~250–5,000 people.
- \$250k+/yr model spend, multi-provider, growing fast.
- Buyer: platform/eng lead + VP Finance / FinOps.

WHAT YOU WALK AWAY WITH — MEASURED ON YOUR OWN DATA

Attribution coverage — % of spend joined to a ticket/feature, with a fidelity tier on every dollar	your number
Forecast accuracy — leave-one-out back-test error on your delivered work	your number
One program budget — a body of work capped, paced, with a projected-breach date	live
Commitment read — recommended on-demand vs. committed-spend posture + estimated savings	advisory

We measure these on your traffic — we don't assert a benchmark. If the attribution gap isn't real for you, we'll say so.

THE PILOT

- Read-only / metadata-only connect — no app rewrite, nothing in your request path.
- ~2 weeks to a measured attribution-coverage % and a forecast back-test.
- Free; converts to an annual scoped to your AI spend under management.
- Optional savings-share later — measured, no clawback. We only bill on proven savings.